

# Forecasting

---

# Forecasting Is Extrapolating!

**First, note that:**

Forecasting is ***extrapolating***

## STAT 5371 Example

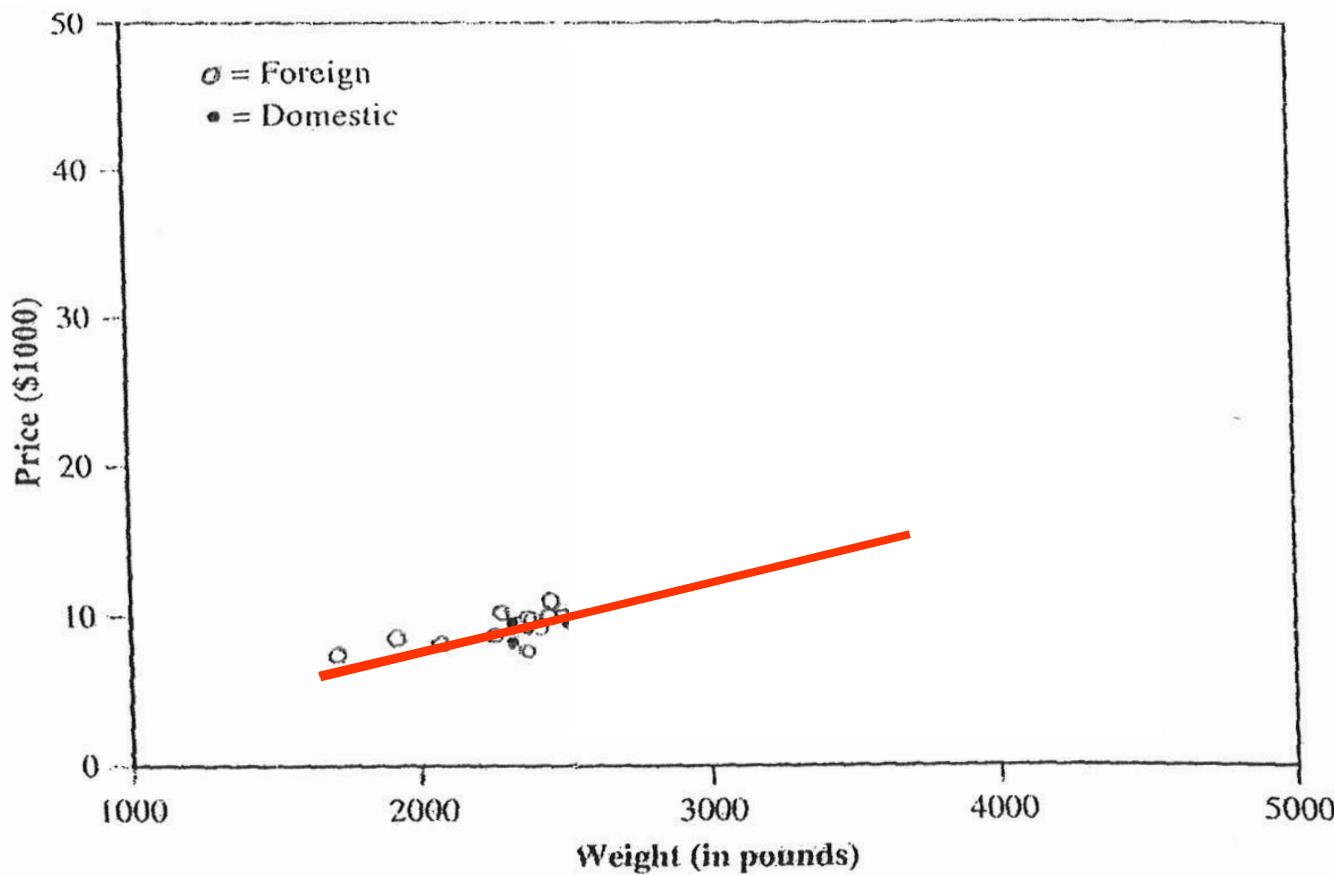


FIGURE 2.9 Base price in thousands of dollars versus weight in pounds for 1991 four-door sedans (Exercise 2.7).

Predict the price of a car that weighs 3500 pounds.

- Extrapolation would say that it's about \$16,000.

## STAT 5371 Example

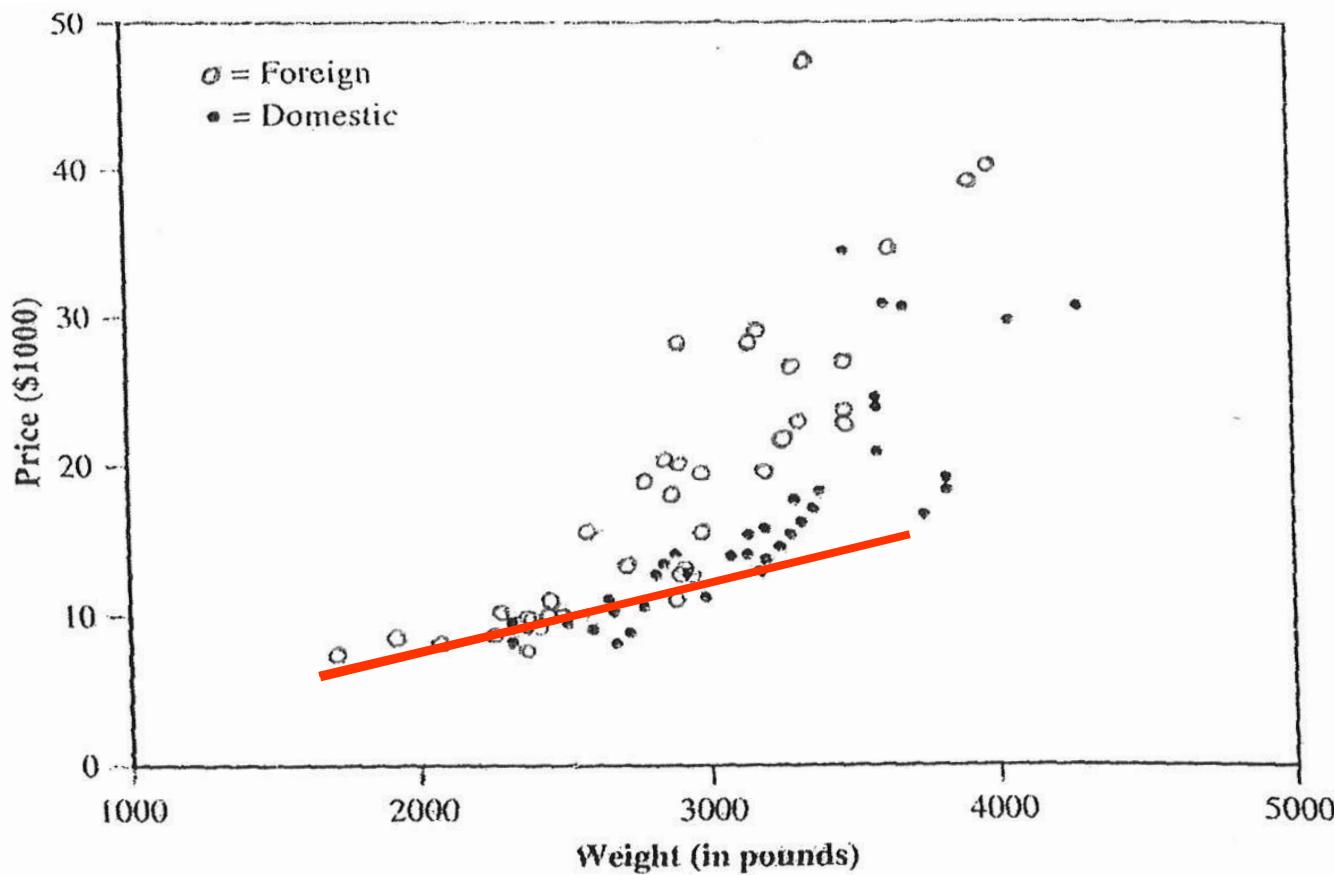


FIGURE 2.9 Base price in thousands of dollars versus weight in pounds for 1991 four-door sedans (Exercise 2.7).

Predict the price of a car that weighs 3500 pounds.

- Extrapolation would say that it's about \$16,000.

**Oops!**

## **Note: Forecasting is extrapolating.**

Lesson from regression analysis and 5371 example?

- ***Be careful when extrapolating!***

However, it's not very important to:

- “Predict” the sunspot number for 2012
- “Predict” sales for the previous quarter

...

Sometimes we need to extrapolate!

That is, sometimes we need to forecast!

**DataScience@SMU**

# Various Forecasts

---

# Deterministic Signal-Plus-Noise Models

$$X_t = s_t + Z_t$$

$s_t$  is a deterministic signal

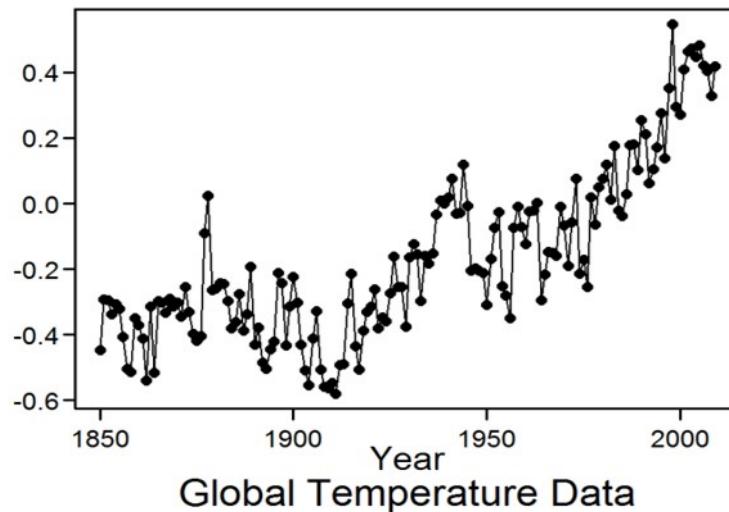
$Z_t$  is a zero-mean, stationary process

**Example signals:**  $s_t = a + bt$

$$s_t = a + bt + ct^2$$

$$s_t = A \cos(2\pi ft + C), C \text{ constant}$$

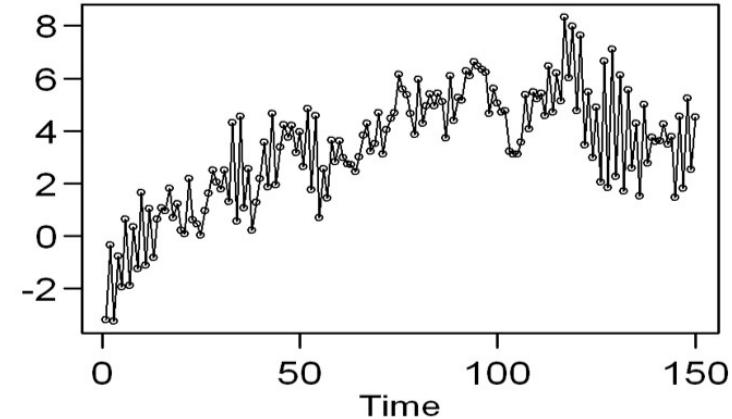
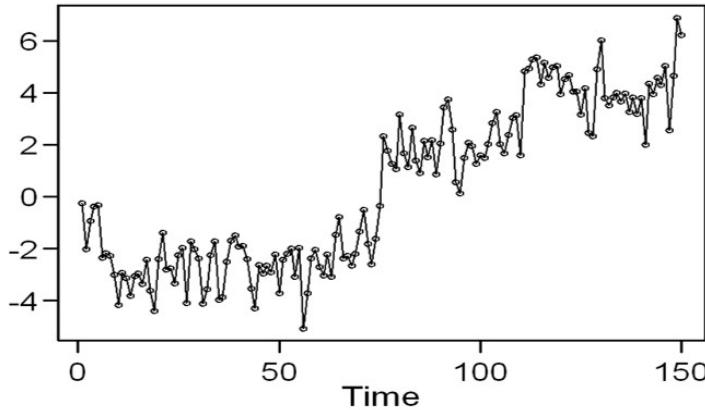
**Recall:** Sometimes it's not easy to tell whether a deterministic signal is present in the data



**Is there a deterministic signal?**

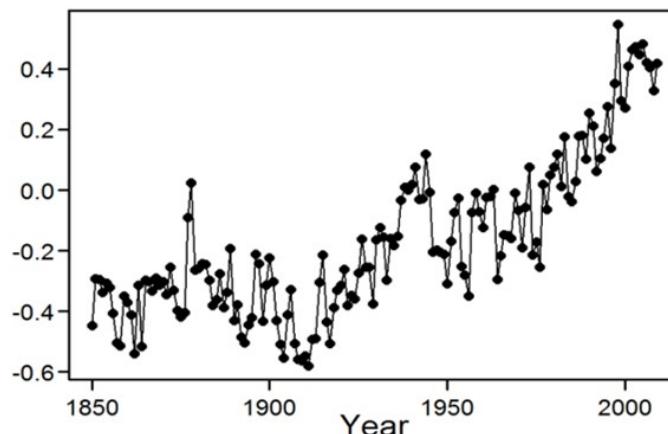
## Realizations:

- Is there a deterministic signal?



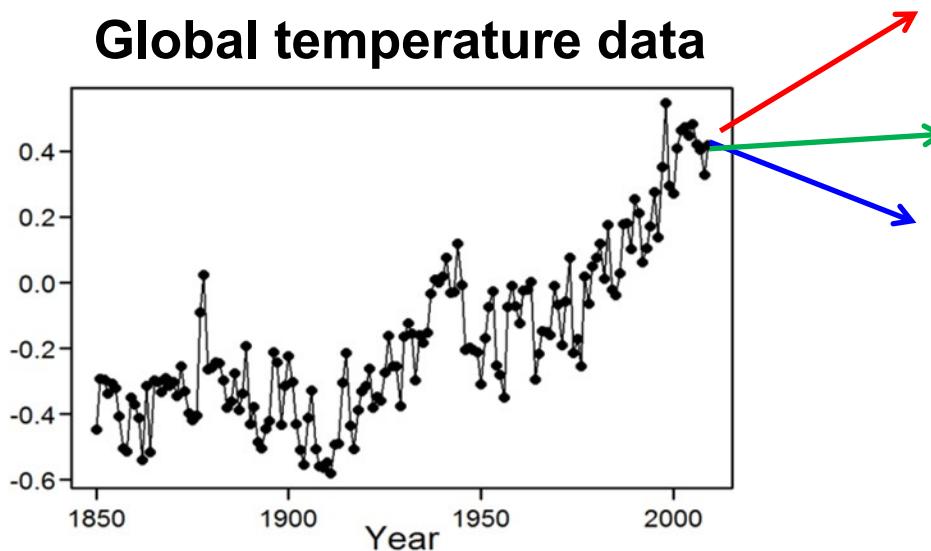
**Recall:** Sometimes it's not easy to tell whether a deterministic signal is present in the data.

Global temperature data



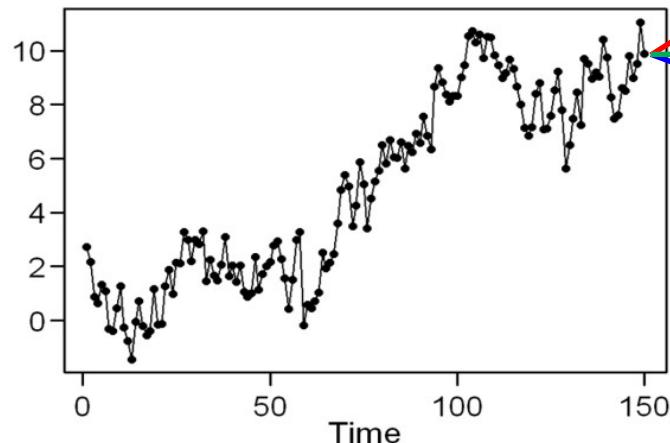
# How Would You Predict/Forecast into the Future?

It depends on the model.

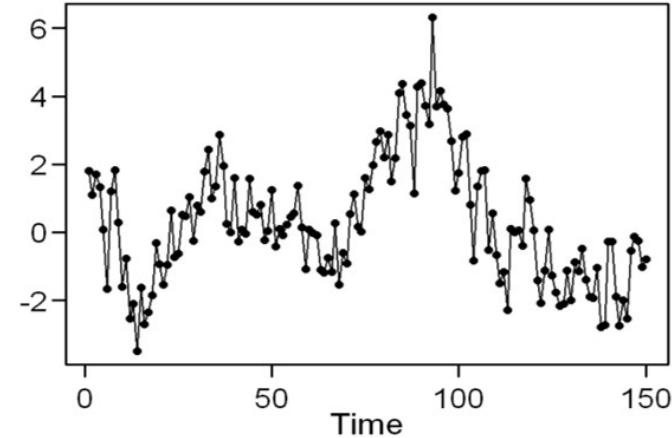
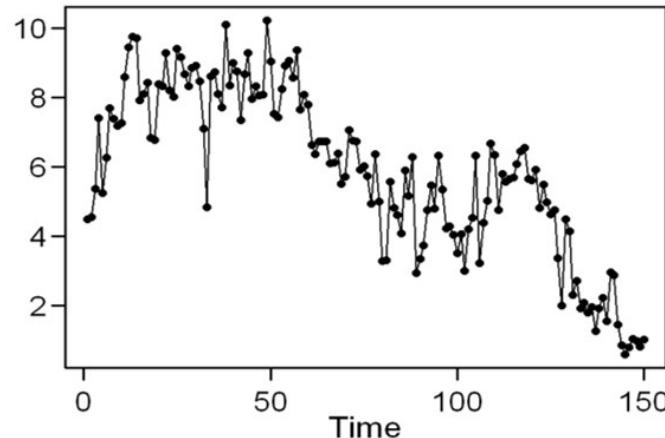
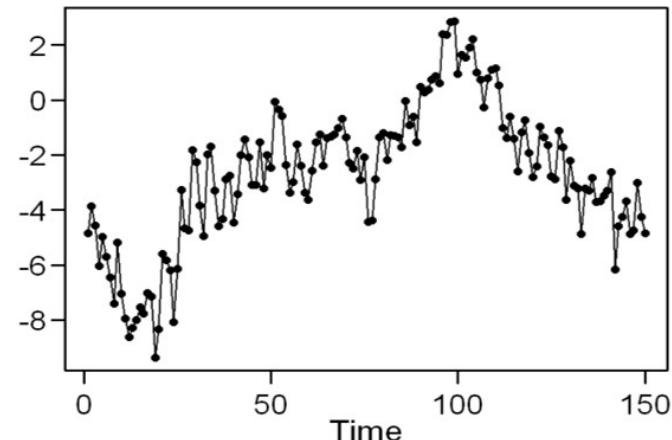


Suppose I told you that this is a realization from the stationary model

$$(1 - .99B)(1 - .2B + .455B^2)(1 - .53B)X_t = a_t$$



Other realizations  
(from this model)



Now which forecast do you choose?

**DataScience@SMU**

# Forecasting Setting

---

# Forecasting Setting in this Chapter

**Forecast future behavior of a time series given a finite realization of its past.**

- The use of ARMA models for this purpose has become popular in recent years.

## ARMA forecasting vs. curve fitting

- Curve fitting (regression)
  - Underlying assumption that future behavior follows some deterministic path with only random fluctuations
- ARMA and other time series-based forecasting
  - Underlying assumption is that the future is guided only by its correlation to the past

# Box-Jenkins Approach

G. E. P. Box and G. M. Jenkins popularized the use of ARMA models for forecasting.

- The most recent version of their classic text is Box, Jenkins, and Reinsel (2008).
- We will utilize their approach for obtaining ARMA, ARIMA, and seasonal forecasts.
  - In lecture and in **tswge**

## Important

In this unit, we will assume that we know the true model (i.e.,  $p$ ,  $q$ , the  $\varphi$ 's and  $\theta$ 's).

**DataScience@SMU**

# Strategy and Notation

---

# Strategy and Notation

Time series data are available:  $X_1, X_2, \dots, X_{t_0}$

Goal is to forecast a future (usually unknown) value, say  $X_{t_0+4}$

## Note:

The typical case is  $t_0 = n$ , i.e.  $X_{t_0} = X_n$  is the last value in the observed realization

- however, we may want to forecast, say, the last 4 values in the realization (which are known values) to assess the performance of our forecasts
  - in this case  $t_0 = n - 4$

# Strategy and Notation

$\hat{X}_{t_0}(\ell)$  is the forecast of  $X_{t_0+\ell}$  given data up to time  $t_0$

- $t_0$  is called the ***forecast origin***
- $\ell$  is the ***lead time***, i.e. the number of time units (***steps ahead***) which we want to forecast

**For example**, suppose we observe  $X_1, X_2, \dots, X_{10}$  ( $= X_{t_0}$ ).

and we want to forecast  $X_{12}$  ( $= X_{t_0+2}$ ).

Then  $\hat{X}_{10}(2)$  is the forecast of  $X_{12}$

And  $\hat{X}_{10}(3)$  is the forecast of  $X_{13}$

And  $\hat{X}_{10}(7)$  is the forecast of  $X_{17}$

And  $\hat{X}_{10}(l)$  is the forecast of  $X_{10+l}$

# Strategy and Notation

$\hat{X}_{t_0}(\ell)$  is the forecast of  $X_{t_0+\ell}$  given data up to time  $t_0$

- $\hat{X}_{t_0}(\ell) = X_{t_0+\ell}$  if  $\ell \leq 0$

**For example**, suppose we observe  $X_1, X_2, \dots, X_{10} (= X_{t_0})$ .

$\hat{X}_{10}(\ell)$  is a forecast of  $X_{10+\ell}$ , so

$\hat{X}_{10}(0)$  is a forecast of  $X_{10+0} = X_{10}$ , which has **already been observed**.

It makes intuitive sense that  $\hat{X}_{10}(0) = X_{10}$

$$\hat{X}_{10}(-1) = X_9$$

$$\hat{X}_{10}(-2) = X_8$$

And so on

DataScience@SMU

# Forecasting with AR(1)

---

# Forecasting Using an AR(1) Model

Suppose that  $X_t$  is AR(1) (i.e.,  $X_t - \varphi_1 X_{t-1} = (1 - \varphi_1)\mu + a_t$ ).

And suppose again we observe  $X_1, X_2, \dots, X_{10}$  ( $= X_{t_0}$ ).

So,  $X_{10+2} = \varphi_1 X_{10+1} + a_{10+2} + \begin{pmatrix} 1 & 1 \end{pmatrix}$  **Light board**

i.e.,  $X_{12} = \varphi_1 X_{11} + a_{12} + \begin{pmatrix} 1 & 1 \end{pmatrix}$

## Notes:

- recall that we are assuming we know  $\varphi_1$
- however, we still can't use this formula to calculate  $X_{12}$  because we don't know:
  - $X_{11}$ ,  $a_{12}$ , and  $\mu$
  - we will estimate  $\mu$  with  $\bar{X}$

# General Forecasting Formula for an AR(1)

$$\hat{X}_{t_0}(\ell) = \varphi_1 \hat{X}_{t_0}(\ell-1) + \bar{X}(1 - \varphi_1)$$

**Light board**

**Note:** This formula is used recursively.

That is, we iteratively calculate:

$$\begin{aligned}\hat{X}_{t_0}(1) \\ \hat{X}_{t_0}(2) \\ \vdots \\ \hat{X}_{t_0}(\ell-1) \\ \boxed{\hat{X}_{t_0}(\ell)}\end{aligned}$$

# Forecasting Using an AR(1) Model (Example)

We observe  $X_1, X_2, \dots, X_{10}$  ( $= X_{t_0}$ ) and want to forecast  $X_{12}$

$$X_{12} = \varphi_1 X_{11} + a_{12} + \bar{X}(1 - \varphi_1)$$

**Light board**

Based on data to time  $t = 10$ , we **still** can't calculate  $X_{12}$  because we don't know  $X_{11}$  or  $a_{12}$ .

**Strategy:** We will need to **forecast  $X_{12}$  iteratively** (based on data to time  $t = 10$ ).

**First:** We forecast  $X_{11}$  based on data to time  $t = 10$ .

$$X_{11} \text{ is given by } X_{11} = \varphi_1 X_{10} + a_{11} + \bar{X}(1 - \varphi_1)$$

replacing actual (unknown) values with forecasts as of  $t = 10$

$$\hat{X}_{10}(1) = \varphi_1 \hat{X}_{10}(0) + \hat{a}_{10}(1) + \bar{X}(1 - \varphi_1)$$

# Forecasting Using an AR(1) Model (Example)

$$\hat{X}_{10}(1) = \varphi_1 \hat{X}_{10}(0) + \hat{a}_{10}(1) + \bar{X}(1 - \varphi_1)$$

**Light board**

**Recall:**  $\hat{X}_{10}(0) = X_{10}$  (known value)

**How about:**  $\hat{a}_{10}(1)$ ?

$a_t$  is random white noise with zero mean, and as of  $t = 10$  we do not know  $a_{11}$  -- but we do know it has zero mean.

So, it makes intuitive sense to define  $\hat{a}_{10}(1) = 0$ , and correspondingly, for  $\ell > 0$ , we define  $\hat{a}_{10}(\ell) = 0$

**Now:**  $\hat{X}_{10}(1) = \varphi_1 X_{10} + \bar{X}(1 - \varphi_1)$  (which we can calculate)

**And finally:** 
$$\begin{aligned}\hat{X}_{10}(2) &= \varphi_1 \hat{X}_{10}(1) + \hat{a}_{10}(2) + \bar{X}(1 - \varphi_1) \\ &= \varphi_1 \hat{X}_{10}(1) + \bar{X}(1 - \varphi_1) \text{ (which again we can calculate)}\end{aligned}$$

# Example: AR(1) forecasts

Suppose that the realization  $X_1, \dots, X_{80}$  is observed from the AR(1) model  $(1 - .8B)(X_t - 25) = a_t$  and further that  $X_{80} = 21.77$  and  $\bar{X} = 24.17$ . (These data are in tswge file fig6.1nf.)

Although we know  $\mu = 25$ , we use  $\bar{X} = 24.17$  since in practice,  $\mu$  will be unknown to us.

Calculate  $\hat{X}_{80}(\ell)$ ,  $\ell = 1, 2, 3$

**Light board**

**Forecast Function :**

$$\hat{X}_{80}(\ell) = .8\hat{X}_{80}(\ell-1) + \bar{X}(1-.8) = .8\hat{X}_{80}(\ell-1) + 24.17(.2)$$

4.83

$$\text{So, } \hat{X}_{80}(1) = .8\hat{X}_{80}(0) + 4.83 = 17.42 + 4.83 = 22.25$$

$$\hat{X}_{80}(2) = .8\hat{X}_{80}(1) + 4.83 = .8(22.25) + 4.83 = 22.63$$

$$\hat{X}_{80}(3) = .8\hat{X}_{80}(2) + 4.83 = .8(22.63) + 4.83 = 22.93$$

**DataScience@SMU**

# Example

---

Forecasting with an AR(1) Model

# Example: AR(1) forecasts

$$\hat{X}_{80}(1) = 22.25$$

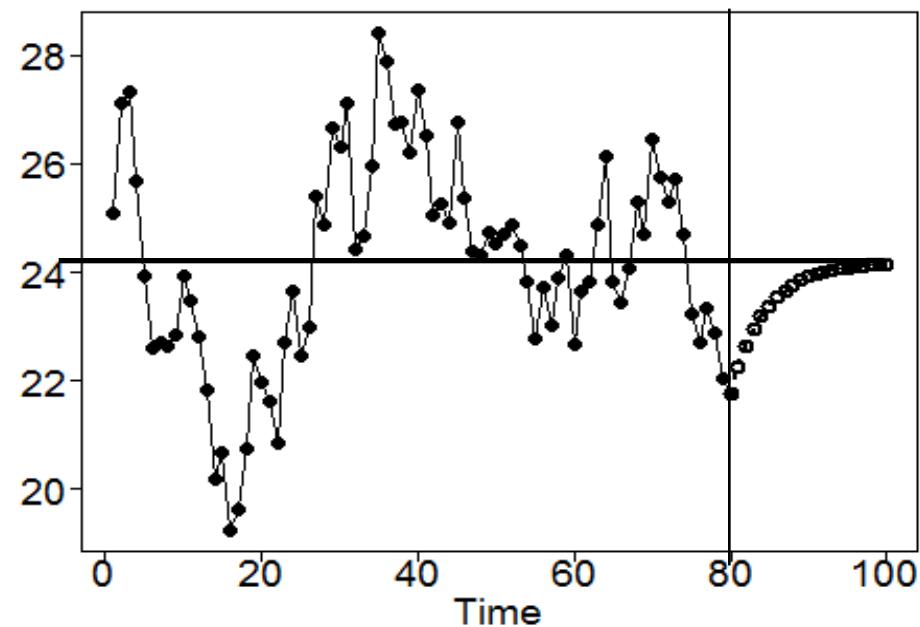
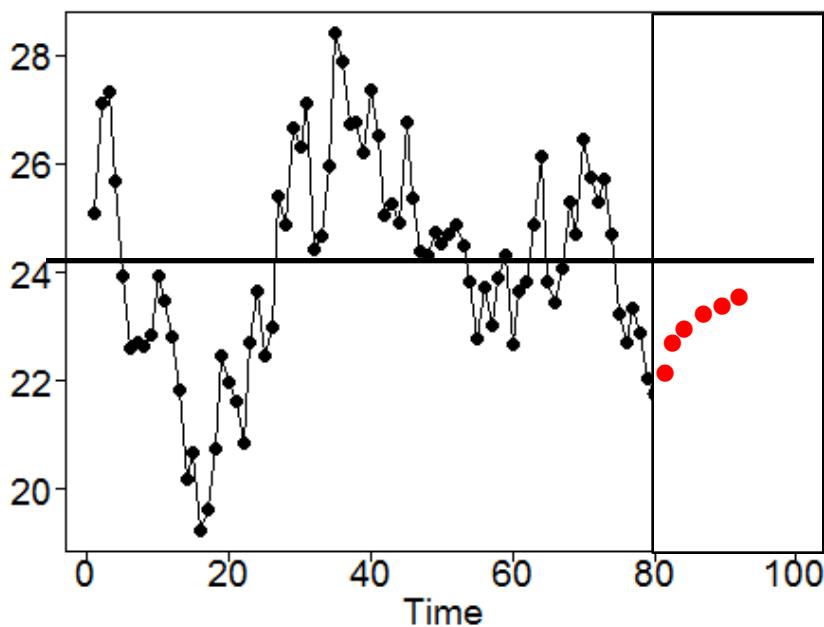
What is  $\hat{X}_{80}(4)$ ? 23.17

$$\hat{X}_{80}(2) = 22.63$$

What is  $\hat{X}_{80}(5)$ ? 23.37

$$\hat{X}_{80}(3) = 22.93$$

What is  $\hat{X}_{80}(6)$ ? 23.52



**DataScience@SMU**

# Forecasting with AR(2) and AR( $p$ )

---

# Basic Formula for AR(p) Forecasts

$$\hat{X}_{t_0}(\ell) = {}_1\hat{X}_{t_0}(\ell-1) + \cdots + {}_p\hat{X}_{t_0}(\ell-p) \\ + \bar{X}(1 \quad 1 \quad \cdots \quad p)$$

## Example: AR(2)

$$\hat{X}_{t_0}(\ell) = \varphi_1 \hat{X}_{t_0}(\ell-1) + \varphi_2 \hat{X}_{t_0}(\ell-2) + \bar{X}(1 - \varphi_1 - \varphi_2)$$

$X_1, X_2, \dots, X_{74}, X_{75}$  is a realization from the AR(2) model

$(1 - 1.6B + .8B^2)(X_t - 30)$  where  $\bar{X} = 29.4$ ,  $X_{74} = 27.7$ ,  $X_{75} = 23.4$

## Forecasts

## Light board

$$\hat{X}_{75}(0) = 1.6\hat{X}_{75}(0) - .8\hat{X}_{75}(-1) + 29.4(1 - 1.6 + .8)$$

$$\hat{X}_{75}(1) = 1.6\hat{X}_{75}(0) - .8\hat{X}_{75}(-1) + 5.9$$

$$\hat{X}_{75}(1) = 1.6X_{75} - .8X_{74} + 5.9 = 1.6(23.4) - .8(27.7) + 5.9 = 21.2$$

$$\hat{X}_{75}(2) = 1.6\hat{X}_{75}(1) - .8X_{75} + 5.8 = 1.6(21.2) - .8(23.4) + 5.9 = 21.1$$

$$\hat{X}_{75}(3) = 1.6(21.1) - .8(21.2) + 5.9 = 22.7$$

**DataScience@SMU**

# Example

---

Forecasting with an AR(2) Model

# Example: AR(2) forecasts

$$\hat{X}_{75}(1) = 21.2$$

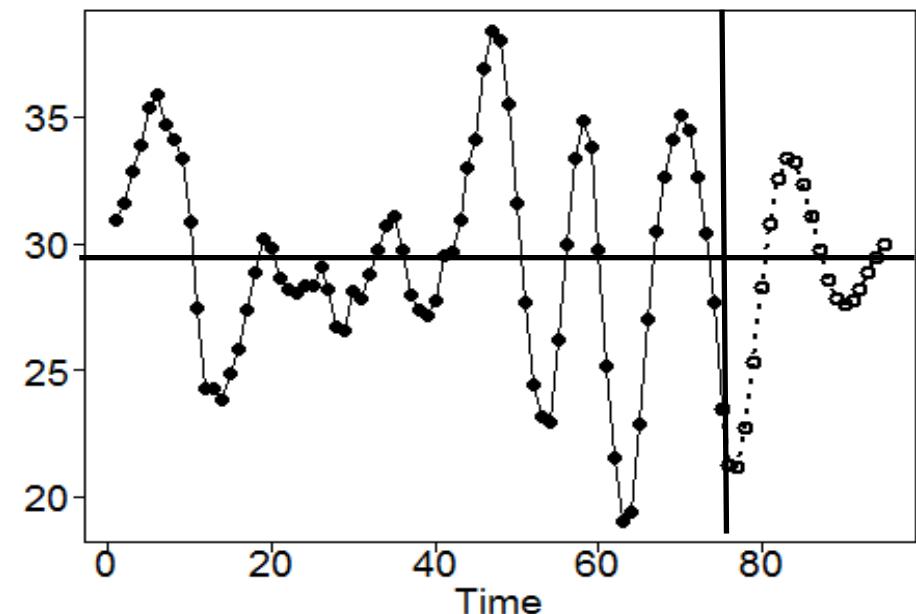
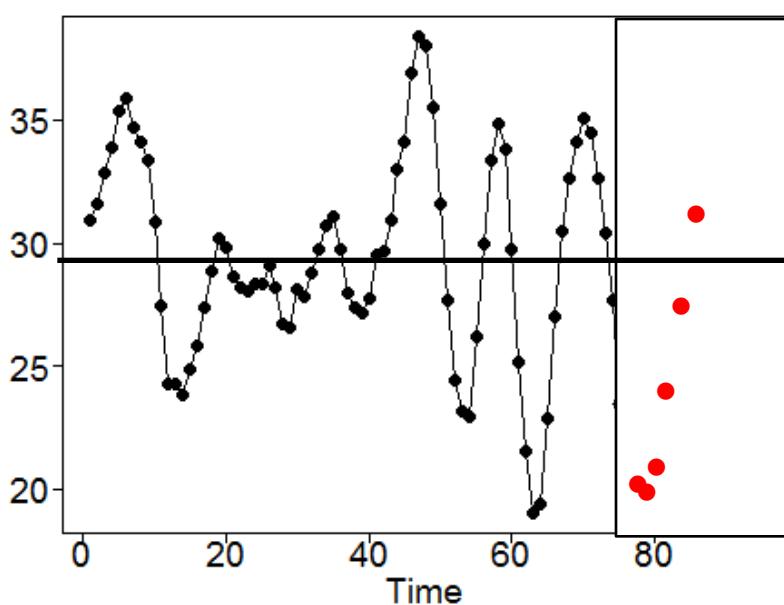
What is  $\hat{X}_{75}(4)$ ? 25.3

$$\hat{X}_{75}(2) = 21.1$$

What is  $\hat{X}_{75}(5)$ ? 28.2

$$\hat{X}_{75}(3) = 22.7$$

What is  $\hat{X}_{75}(6)$ ? 30.8



DataScience@SMU

tswge

---

## Forecasting AR( $p$ ) Models

# tswge Forecasting Examples

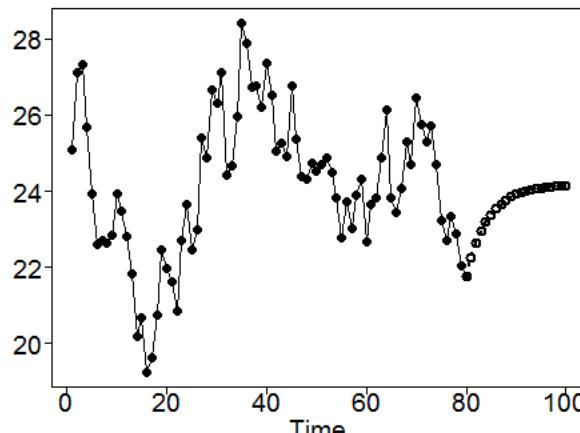
## AR(1) example (revisited)

Realization  $X_1, \dots, X_{80}$  in tswge file fig6.1nf from the AR(1) model

$$(1 - .8B)(X_t - 25) = a_t$$

```
data(fig6.1nf)
fore.arma.wge(fig6.1nf,phi=.8,n.ahead=20,plot=TRUE,limits=FALSE)
```

These commands plot the realization and the n.ahead=20 forecasts shown below. The values of the forecasts are in \$f of the output.



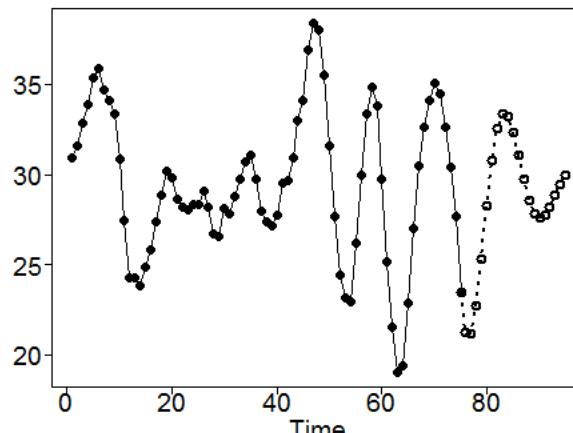
# tswge Forecasting Examples

## AR(2) example (revisited)

Realization  $X_1, \dots, X_{75}$  from the AR(2) model  $(1 - 1.6B + .8B^2)(X_t - 25) = a_t$

```
x2=gen.arma.wge(n=75,phi=c(1.6,-.8),sn=24)
x2=x2+25
fore.arma.wge(x2,phi=c(1.6,-.8),n.ahead=20,limits=FALSE)
```

These commands plot the realization and the `n.ahead=20` forecasts shown below. The values of the forecasts are in \$f of the output.



DataScience@SMU

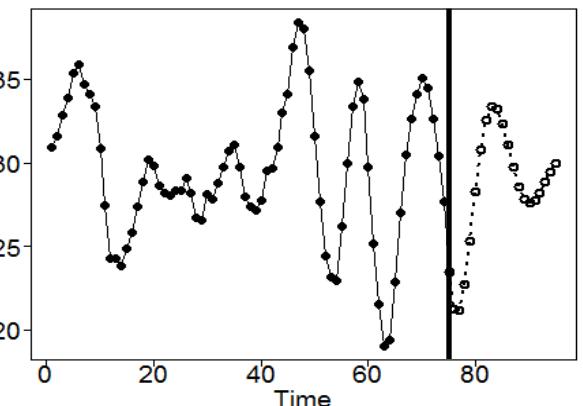
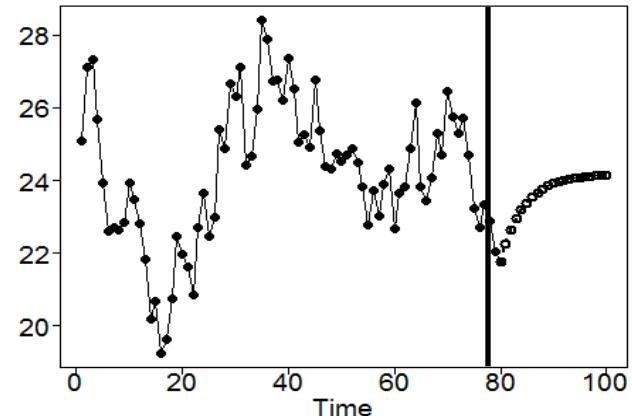
# Eventual Forecast Function

---

# Eventual Forecast Function

Two key observations from two previous forecast plots obtained using:

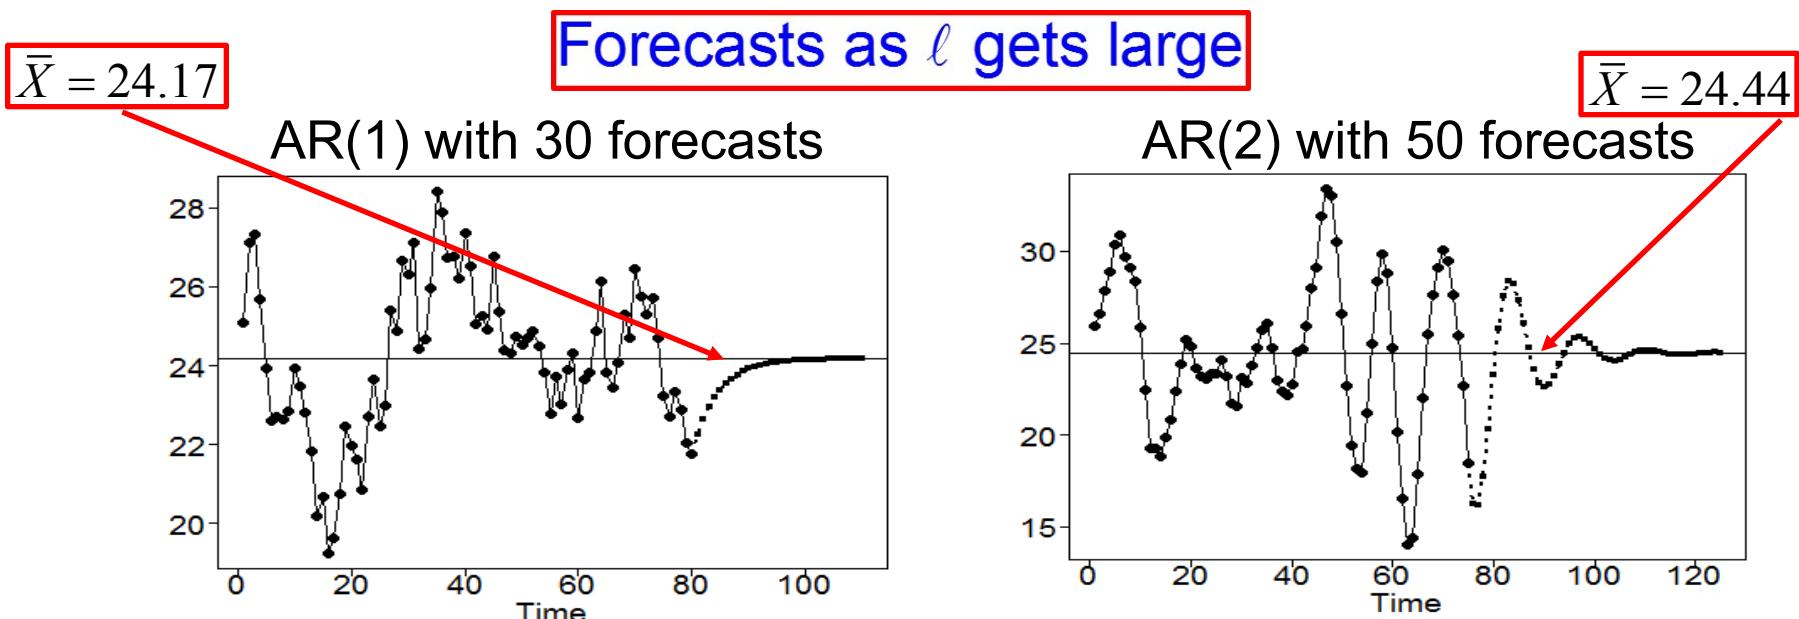
1. In both cases, the forecasts behave like the autocorrelations.
  - AR(1) case: forecasts have exponential shape (although increasing instead of decreasing)
  - AR(2) case: forecasts look like a damped sinusoid (AR part has complex conjugate roots)



# Eventual Forecast Function

2. The forecasts *tend toward the realization mean* as number of steps ahead increases.

- The correlation between observed values (up to time  $t_0$ ) and  $X_{t_0+\ell}$  gets small as  $\ell$  gets large
- Since the observed values give very little information about  $X_{t_0+\ell}$ , so it is reasonable to simply forecast  $\bar{X}$



DataScience@SMU

# Basic Formula for Forecasting with ARMA( $p,q$ )

---

# Basic Formula for ARMA( $p,q$ ) Forecasts

$$\hat{X}_{t_0}(\ell) = \sum_{j=1}^p \varphi_j \hat{X}_{t_0}(\ell-j) - \sum_{j=\ell}^q \theta_j \hat{a}_{t_0+\ell-j} + \bar{X} \left[ 1 - \sum_{i=1}^p \varphi_i \right]$$

**Note:** These forecasts depend on white-noise variance estimates. Obtaining these estimates is computationally intensive. (See Woodward et al. (2017).)

- The MA part of the model affects the forecasts, and the eventual forecasts behave in much the same way as they would for the AR part of the model.

***In practice*,** we will obtain ARMA( $p,q$ ) (and AR( $p$ )) forecasts using **tswge** function **fore.arma.wge**.

**DataScience@SMU**

# tswge

---

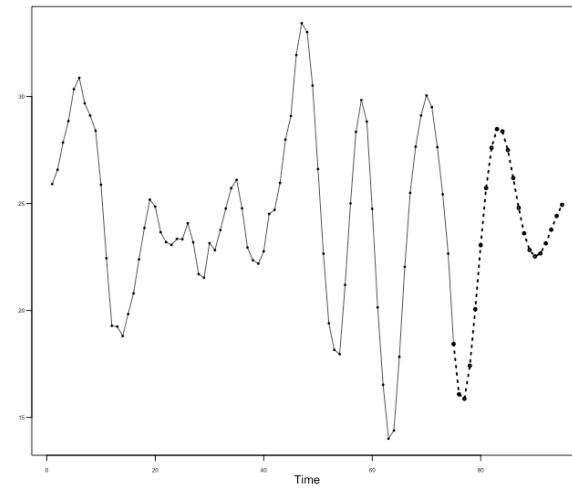
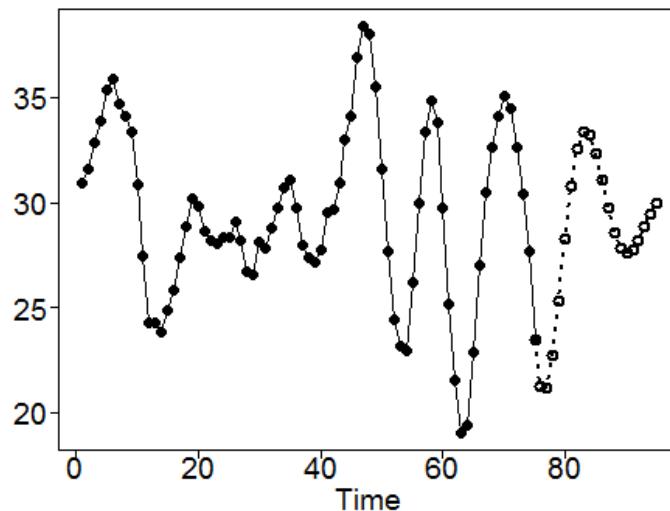
## ARMA( $p,q$ ) Forecast Examples

# tswge Forecasting Examples

ARMA(2,1) example in Woodward et al. (2017)

```
x2=gen.arma.wge(n=75,phi=c(1.6,-.8),sn=24)
x2=x2+25
fore.arma.wge(x2,phi=c(1.6,-.8), n.ahead=20,limits=FALSE)
```

```
fore.arma.wge(x2,phi=c(1.6,-.8), theta = -.9, n.ahead=20,limits=FALSE)
```



See Woodward et al. (2017) Appendix 6B for more details.

DataScience@SMU

# tswge

---

## Canadian Lynx Example

# tswge Forecasting Examples: Canadian Lynx

Show how increasing the order of the AR factor helps grab the cycle in the data AR(4), ARMA(4,1), and AR(11).

$$(1 - 1.3B + 0.7B^2 - 0.1B^3 + 0.2B^4)(X_t - 2.9) = a_t$$

AR(4): phi = c(1.3, -0.7, 0.1, -0.2)

$$(1 - 0.7B - 0.1B^2 + 0.2B^3 + 0.3B^4)(X_t - 2.9) = (1 + .6B)a_t$$

ARMA(4,1): phi = c(0.7, 0.1, -0.2, -0.3), theta = -0.6

AR(11)

phi = c(1.1676, -0.5446, 0.2662, -0.3094, 0.1540, -0.1463, 0.0569, -0.0294, 0.1346, 0.2021, -0.3394)

DataScience@SMU

# Probability Limits and ARMA Forecasts

---

Setting

# Probability Limits for ARMA Forecasts

## Setting

Given  $\hat{X}_{t_0}(\ell)$ , an estimate of  $X_{t_0+\ell}$  based on data to time  $t_0$ , we want to construct limits such that the probability that  $X_{t_0+\ell}$  will fall within the limits is 95%.

## Recall

An ARMA(p,q) model can be expressed as a General Linear Process (GLP) given by

$$X_t = \sum_{j=0}^q a_{t-j} \text{ where } a_0 = 1$$

and  $a_t$  is normal white noise (mean 0 and variance  $\sigma_a^2$ )

# Probability Limits for ARMA Forecasts

## Facts

A.  $e_{t_0}(\ell) = X_{t_0+\ell} - \hat{X}_{t_0}(\ell)$  (forecast error)

$$= \sum_{k=0}^{\ell-1} \psi_k a_{t_0+\ell-k}$$

B.  $e_{t_0}(\ell)$  is normally distributed (linear combination of normals)

C.  $E[e_{t_0}(\ell)] = 0$  (i.e.  $E[X_{t_0+\ell} - \hat{X}_{t_0}(\ell)] = 0$ )

D.  $\text{var}(e_{t_0}(\ell)) = \text{var}\left(\sum_{j=0}^{\ell-1} \psi_j a_{t_0+\ell-j}\right)$

$$= \sigma_a^2 \sum_{j=0}^{\ell-1} \psi_j^2$$

DataScience@SMU

# Psi Weights and ARMA Models

---

# A Note About the $\psi_j$ 's for ARMA( $p,q$ ) Models

$(B)X_t = (B)a_t$  can be expressed as

$$X_t = {}^1(B) (B)a_t$$

$$= \frac{(B)}{(B)} a_t$$

$$= \left( \sum_{j=0}^q B^j \right) a_t$$

**Recall?:** The ratio of two polynomials is a (possibly infinite order) polynomial.

See Woodward et al. (2017) for more details concerning calculating  $\psi$ -weights.

**DataScience@SMU**

# Calculating Psi Weights for ARMA Models

---

**Example:**  $(1 - 1.2B + .6B^2) X_t = (1 - .5B)a_t$

can be written  $X_t = \frac{1 - .5B}{1 - 1.2B + .6B^2} a_t$

Dividing polynomials (remember?)

$$\begin{array}{r} 1 - .5B \\ \hline 1 - 1.2B + .6B^2 \end{array}$$

**Light board**

$$\begin{array}{r} 1 + .7B + .24B^2 - .132B^3 - \dots \\ \hline 1 - .5B \\ \hline 1 - 1.2B + .6B^2 \\ \hline .7B - .6B^2 \\ \hline .7B - .84B^2 + .42B^3 \\ \hline .24B^2 - .42B^3 \\ \hline .24B^2 - .288B^3 + .144B^4 \\ \hline -.132B^3 - .144B^4 \end{array}$$

**Example:**  $(1 - 1.2B + .6B^2) X_t = (1 - .5B)a_t$

can be written  $X_t = \frac{1 - .5B}{1 - 1.2B + .6B^2} a_t$

So, in this case,  $\psi_0 = 1$ ,  $\psi_1 = .7$ ,  $\psi_2 = .24$ ,  $\psi_3 = -.132$ , ...

**Light board**

**DataScience@SMU**

# tswge and Psi Weights

---

**In practice:** We will find  $\psi$ -weights using tswge function psi.weights.wge.

**Return to example:**  $(1 - 1.2B + .6B^2) X_t = (1 - .5B)a_t$ .  
We calculated the  $\psi$ -weights to be  $\psi_0 = 1$ ,  $\psi_1 = .7$ ,  $\psi_2 = .24$ ,  $\psi_3 = -.132$ , ...

### Using tswge:

```
psi.weights.wge(phi=c(1.2,-.6),theta=.5,lag.max=5)
```

gives the output (for  $\psi_1, \dots, \psi_5$ )

```
[1] 0.70000 0.24000 -0.13200 -0.30240 -0.28368
```

which agree with the above calculations.

**Note:**  $\psi_0$  is always equal to 1.

**DataScience@SMU**

# Probability Limits for ARMA Forecasts

---

# Probability Limits for ARMA Forecasts

$$\frac{e_{t_0}(\ell) - 0}{\sigma_a \left\{ \sum_{j=0}^{\ell-1} \psi_j^2 \right\}^{1/2}} \sim N(0,1)$$

Light board

i.e.  $\frac{X_{t_0+\ell} - \hat{X}_{t_0}(\ell)}{\sigma_a \left\{ \sum_{j=0}^{\ell-1} \psi_j^2 \right\}^{1/2}} \sim N(0,1)$

$$\text{Prob} \left[ \hat{X}_{t_0}(\ell) - z_{\alpha/2} \sigma_a \left\{ \sum_{j=0}^{\ell-1} \psi_j^2 \right\}^{1/2} \leq X_{t_0+\ell} \leq \hat{X}_{t_0}(\ell) + z_{\alpha/2} \sigma_a \left\{ \sum_{j=0}^{\ell-1} \psi_j^2 \right\}^{1/2} \right] = .95$$

**DataScience@SMU**

# 95% Probability Limits for Forecasts

$$\hat{X}_{t_0}(\ell) \pm z_{\alpha/2} \hat{\sigma}_a \left\{ \sum_{j=0}^{\ell-1} \psi_j^2 \right\}^{1/2}$$

## Notes:

- (1) 95% applies to individual forecasts
- (2) The residuals,  $a_t$ , are unknown and must be estimated.
  - tswge uses "backcasting" to find the  $\hat{a}_t$ 's.
  - $\hat{\sigma}_a^2$  is the sample variance of the  $\hat{a}_t$ 's.
  - See Woodward et al. (2017) Chapters 6 and 7.

# Probability Limits for ARMA Forecasts with an Example

---

**ARMA(2,1) example:**  $(1 - 1.2B + .6B^2)(X_t - 50) = (1 - .5B)a_t$

$$\hat{X}_{t_0}(\ell) \pm 1.96 \hat{\sigma}_a \left\{ \sum_{j=0}^{\ell-1} \psi_j^2 \right\}^{1/2}$$

Estimate  $\hat{a}$  using  $\hat{a} = .8659$   
 -weights:  $w_0 = 1, w_1 = .7, w_2 = .24, w_3 = .132$

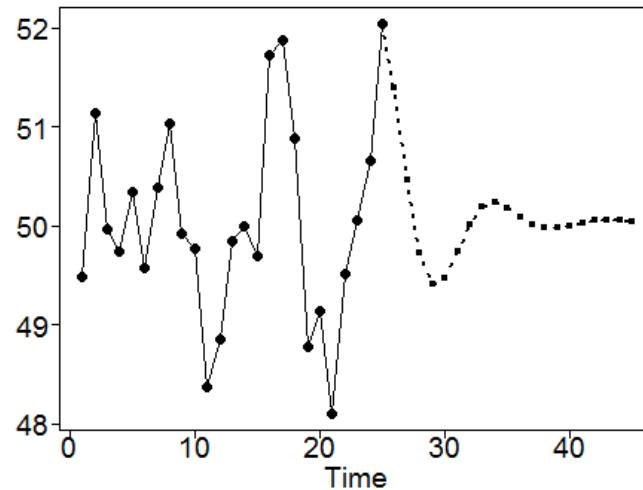
$$\ell = 1 \quad 1.96(.8659)(1) = 1.70$$

$$\ell = 2 \quad 1.96(.8659)\sqrt{1 + .7^2} = 2.07$$

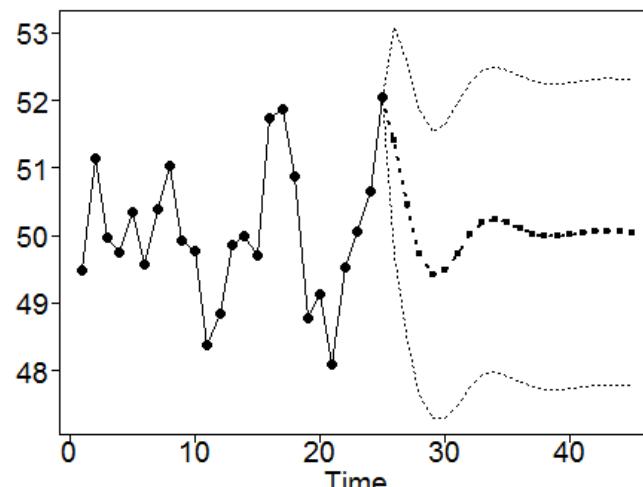
$$\ell = 3 \quad 1.96(.8659)\sqrt{1 + .7^2 + .24^2} = 2.11$$

| $\ell$ | Forecast | Half-width |
|--------|----------|------------|
| 1      | 51.40    | 1.70       |
| 2      | 50.46    | 2.07       |
| 3      | 49.73    | 2.11       |
| 4      | 49.42    | 2.12       |

**ARMA(2,1) example:**  $(1 - 1.2B + .6B^2)(X_t - 50) = (1 - .5B)a_t$



**Forecasts with 95% probability limits**



DataScience@SMU

# tswge and Probability Limits (aka Forecast Limits)

---

# tswge demo: showing forecast limits

To forecast from a stationary ARMA model and show 95% probability limits, use

```
fore.arma.wge(x,phi,theta,n.ahead,plot=TRUE,  
limits=TRUE)
```

**Example:** Forecasts using the AR(1) model  $(1 - .8B)(X_t - 25) = a_t$

```
data(fig6.1nf)
```

```
fore.arma.wge(fig6.1nf,phi=.8,n.ahead=20,  
plot=TRUE,limits=FALSE)
```

```
fore.arma.wge(fig6.1nf,phi=.8,n.ahead=20,  
plot=TRUE,limits=TRUE)
```

# tswge demo

**ARMA(2,1) example:**  $(1 - 1.2B + .6B^2)(X_t - 50) = (1 - .5B)a_t$

data (fig6.2nf)

```
fore.arma.wge(fig6.2nf, phi=c(1.2, -.6), theta=.5,  
n.ahead=20, plot=TRUE, limits=FALSE)
```

```
fore.arma.wge(fig6.2nf, phi=c(1.2, -.6), theta=.5,  
n.ahead=20, plot=TRUE, limits=TRUE)
```

**DataScience@SMU**

# Checking Forecasts

---

ASE

# Checking Forecasts

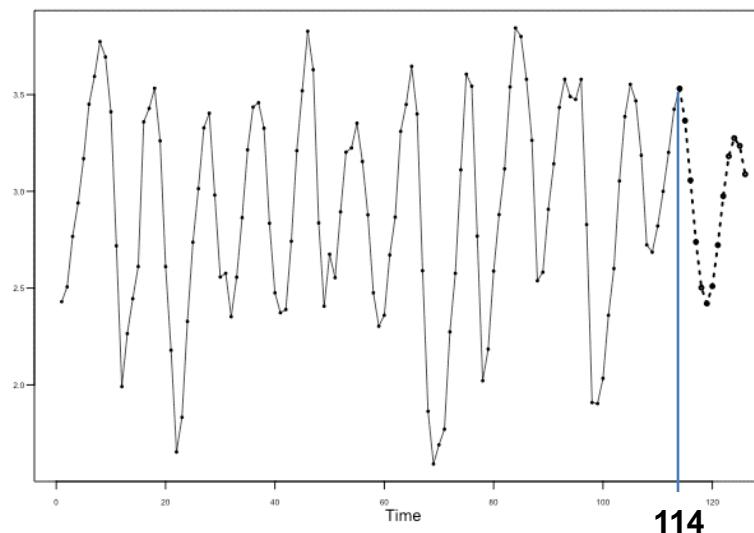
If  $n$  is the length of your time series realization, then sometimes it is useful to set  $t_0 = n - k$  and “forecast the last  $k$  values in the realization.”

- In this way, you can “check” your forecasts with the actual values.
- `tswge` function allows you to do this using the “`lastn`” option.
  - If `lastn=F`, then you forecast past the end of the realization as we have in the previous examples (`lastn=F` is the default option).
  - If `lastn=T`, then you forecast the last `n.ahead` values in the series.

# Checking Forecasts

Again, consider the log lynx dataset. Previously, we have looked at an AR(4) model for these data:

$$(1 - 1.3B + 0.7B^2 - 0.1B^3 + 0.2B^4)(X_t - 2.9) = a_t$$



```
f = fore.arma.wge(llynx,phi = c(1.3, -0.7, 0.1, -0.2), n.ahead = 12, limits = FALSE)
```

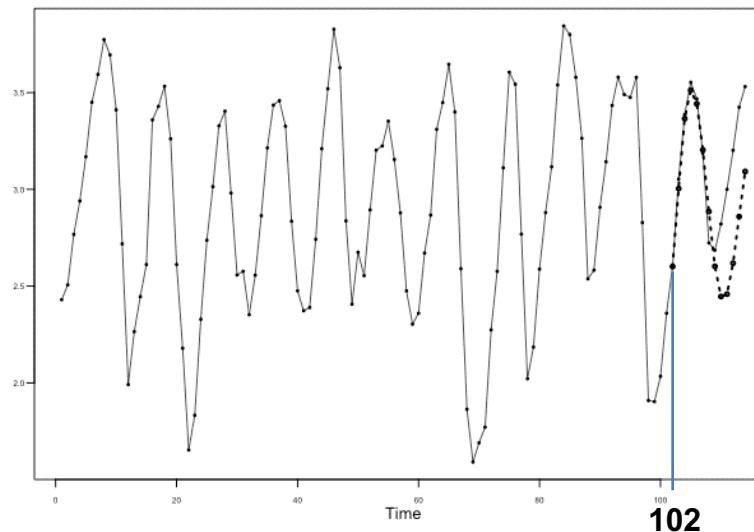
But how good are these forecasts?

We can't know because we don't know the actual values.

# Checking Forecasts

Again, consider the log lynx dataset. Previously, we have looked at an AR(4) model for these data:

$$(1 - 1.3B + 0.7B^2 - 0.1B^3 + 0.2B^4)(X_t - 2.9) = a_t$$



```
f = fore.arma.wge(llynx,phi = c(1.3, -0.7, 0.1, -0.2), n.ahead = 12, lastn = TRUE, limits = FALSE)
```

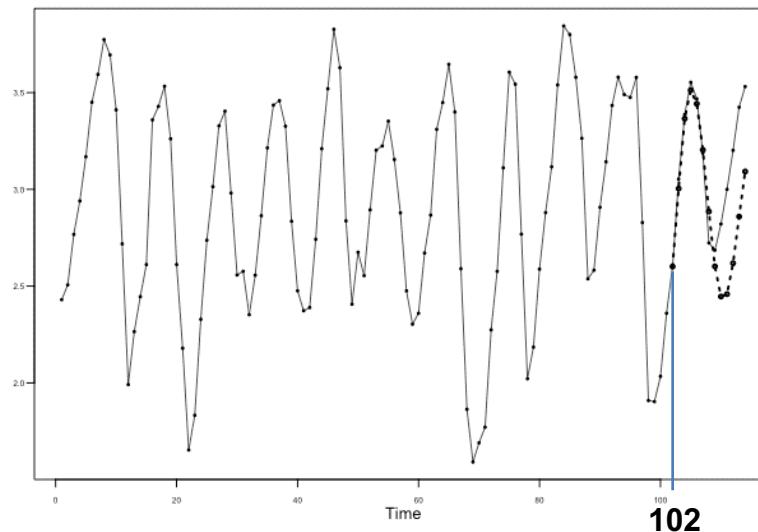
We can forecast the last  $n$  (12 in this case) values of the series so that we can evaluate the fit.

We will use the average square error (ASE):  $ASE = \frac{\sum(\hat{X}_i - X_i)^2}{n}$

# Checking Forecasts

Again, consider the log lynx dataset. Previously, we have looked at an AR(4) model for these data:

$$(1 - 1.3B + 0.7B^2 - 0.1B^3 + 0.2B^4)(X_t - 2.9) = a_t$$



```
f = fore.arma.wge(llynx, phi = c(1.3, -0.7, 0.1, -0.2), n.ahead = 12, lastn = TRUE, limits = FALSE)
```

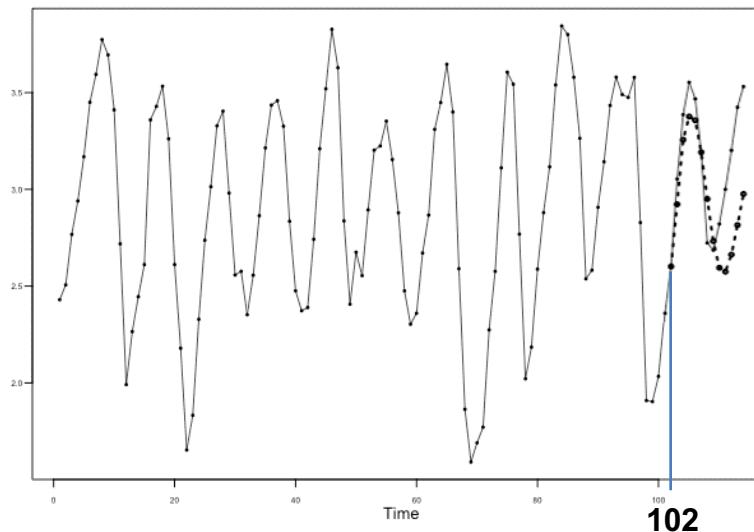
$$ASE = \frac{\sum(\hat{X}_i - X_i)^2}{n} = \frac{\sum_{i=1}^{12}(\hat{X}_i - X_i)^2}{12} \rightarrow > ASE = \text{mean}((f\$f - llynx[103:114])^2)$$

```
> ASE  
[1] 0.1102716
```

# Checking Forecasts

Next, consider the ARMA(4,1) model we looked at for this data set. Let's again calculate the ASE for the last 12 observations.

$$1 - 0.7B - 0.1B^2 + 0.2B^3 + 0.3B^4)(X_t - 2.9) = (1 + .6B)a_t$$



```
f = fore.arma.wge(llynx, phi = c(0.7, 0.1, - 0.2, - 0.3), theta = -.6, n.ahead = 12, lastn = TRUE, limits = FALSE)
```

$$ASE = \frac{\sum(\hat{X}_i - X_i)^2}{n} \rightarrow > \text{ASE} = \text{mean}((f\$f - llynx[103:114])^2)$$

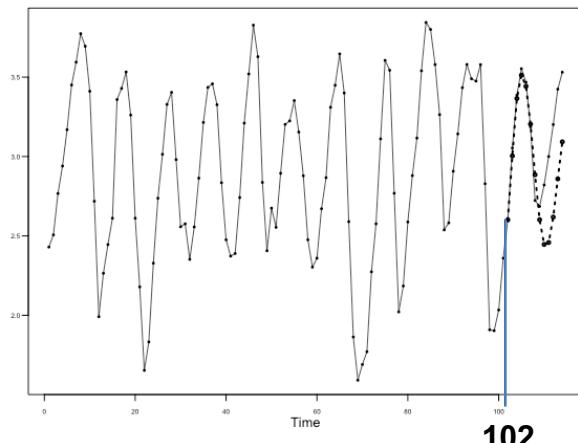
> ASE  
[1] 0.1109845

# Checking Forecasts

Next, consider the ARMA(4,1) model we looked at for this data set. Let's again calculate the ASE for the last 12 observations.

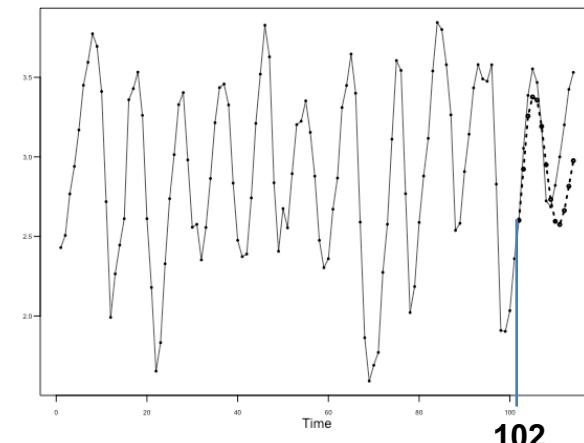
AR(4)

$$(1 - 1.3B + 0.7B^2 - 0.1B^3 + 0.2B^4)(X_t - 2.9) = a_t$$



ARMA(4,1)

$$1 - 0.7B - 0.1B^2 + 0.2B^3 + 0.3B^4)(X_t - 2.9) = (1 + .6B)a_t$$



```
> ASE = mean((f$f-lynx[103:114])^2)  
> ASE  
[1] 0.1102716
```

```
> ASE = mean((f$f-lynx[103:114])^2)  
> ASE  
[1] 0.1109845
```

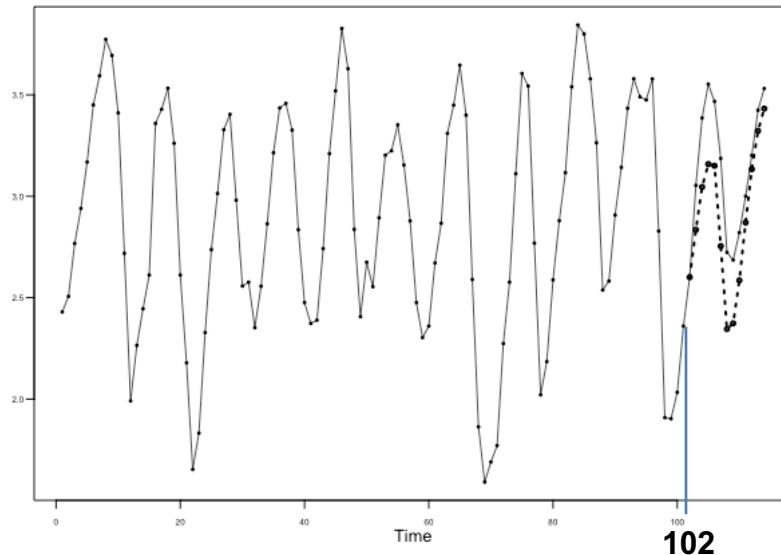
The ASE for the AR(4) model is smaller; therefore, in terms of ASE, the evidence suggests the AR(4) model fits better than the AR(4,1).

# Example: Canadian Lynx Data

A classic model fit to the log lynx data is an AR(11) (Tong, 1977).

```
data(llynx)
f = fore.arma.wge(llynx,phi=c(1.17, -0.54, 0.27, -0.31, 0.15, -0.15,
0.06, -0.03, 0.13, 0.20, -0.34),n.ahead=12,limits=FALSE, lastn = TRUE)
```

## Forecasting last 12 values

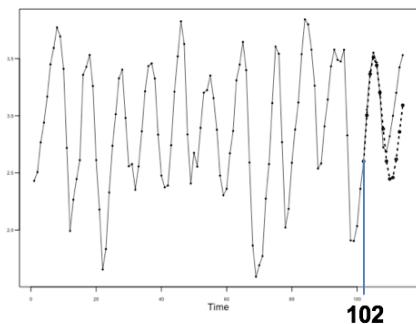


```
mean((f$f-llynx[114-12+1])^2
[1] 0.07865787
```

# Example: Canadian Lynx Data

AR(4)

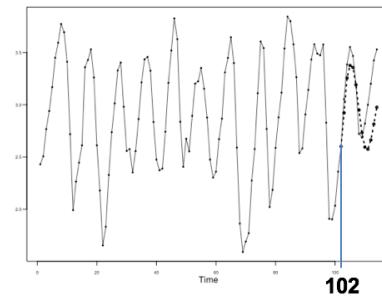
$$(1 - 1.3B + 0.7B^2 - 0.1B^3 + 0.2B^4)(X_t - 2.9) = a_t$$



```
> ASE = mean((f$f-llynx[103:114])^2)  
> ASE  
[1] 0.1102716
```

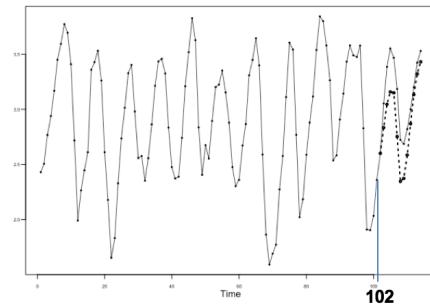
ARMA(4,1)

$$1 - 0.7B - 0.1B^2 + 0.2B^3 + 0.3B^4)(X_t - 2.9) = (1 + .6B)a_t$$



```
> ASE = mean((f$f-llynx[103:114])^2)  
> ASE  
[1] 0.1109845
```

Forecasting last 12 values from AR(11)



```
mean ((f$f-llynx[114-12+1])^2  
[1] 0.07865787
```

**DataScience@SMU**

tswge

---

## Checking Forecasts

## tswge demo: “forecasting” last n.ahead values in your realization

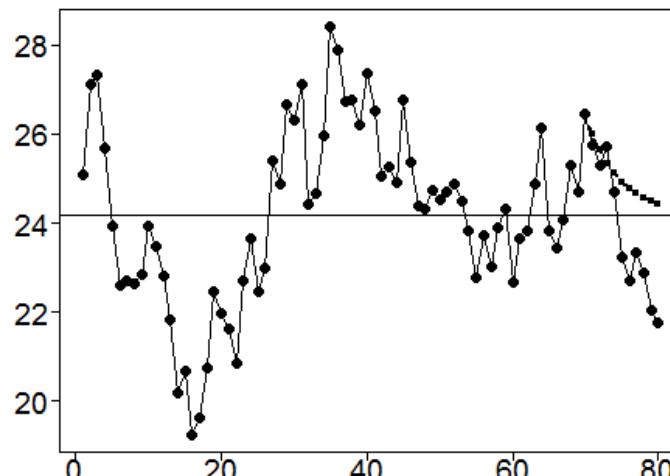
To forecast the last n.ahead values in your data set, use

```
fore.arma.wge(x,phi,theta,n.ahead,lastn=T,  
plot=T,limits=F)
```

**Example:** Forecasts using the AR(1) model  $(1 - .8B)(X_t - 25) = a_t$

```
data(fig6.1nf)
```

```
fore.arma.wge(fig6.1nf,phi=.8,n.ahead=10,lastn=TRUE,  
plot=TRUE,limits=FALSE)
```



## tswge demo: “forecasting” last n.ahead values in your realization

To forecast the last n.ahead values in your data set, use

```
fore.arma.wge(x,phi,theta,n.ahead,lastn=TRUE)
```

**ARMA(2,1) example:**  $(1 - 1.2B + .6B^2)(X_t - 50) = (1 - .5B)a_t$

The ARMA(2,1) data set in fig6.2nf only has 25 data values.

- In order to illustrate the “lastn” feature, we generate a longer realization (n=100) from this model.
- Notice in the code below that we used sn=8 so that this particular realization can be regenerated.
- On the following slide, we illustrate forecasting the last 5, 10, and 20 data values.

# tswge demo: “forecasting” last n.ahead values

**ARMA(2,1) example:**  $(1 - 1.2B + .6B^2)(X_t - 50) = (1 - .5B)a_t$

```
x21=gen.arma.wge(n=100,phi=c(1.2,-.6),theta=.5,sn=8)
```

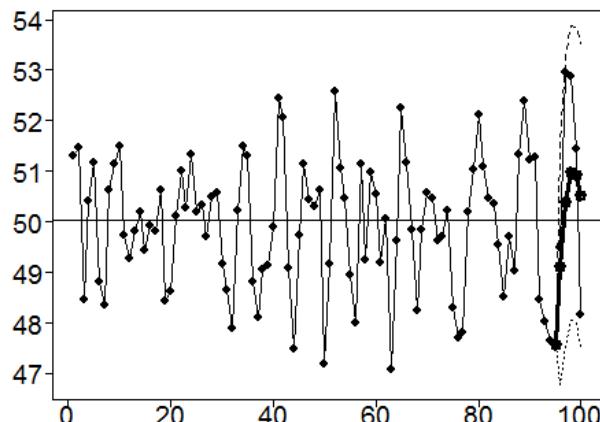
We can recreate this realization.

```
x21=x21+50
```

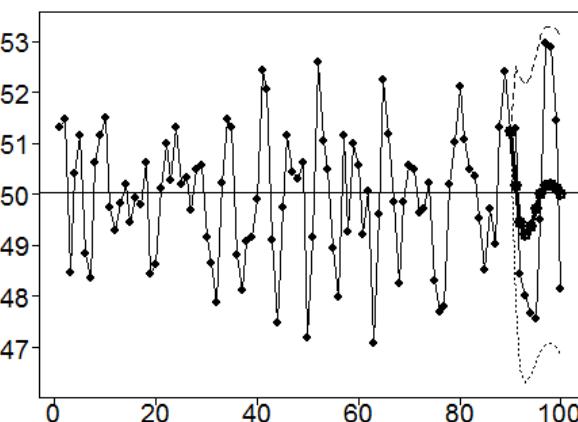
```
fore.arma.wge(x21,phi=c(1.2,-.6),theta=.5,n.ahead=5, lastn=TRUE)
```

We also use 10 and 20.

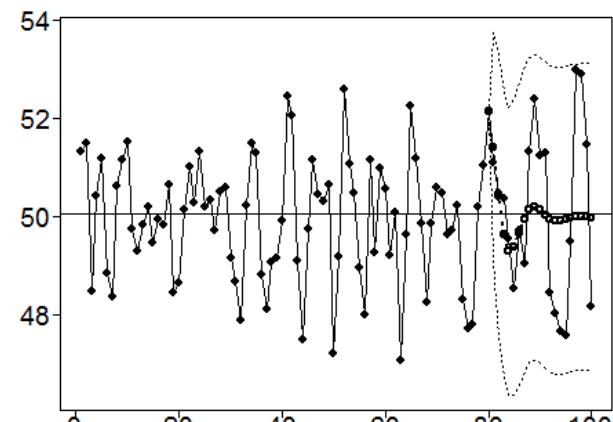
Last 5



Last 10



Last 20



# tswge demo: “forecasting” last n.ahead values

## Notes:

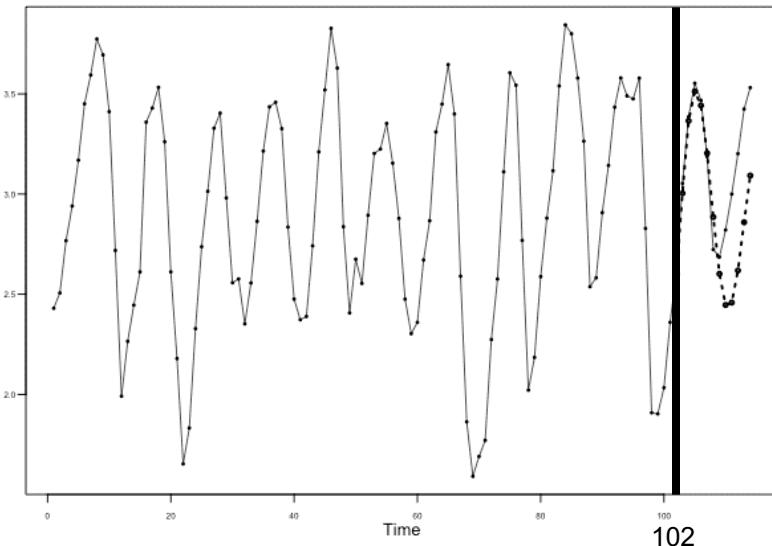
- Short-term forecasts tended to be pretty good.
- For longer steps ahead, the forecasts tended to the mean, 50, shown by the horizontal line.
- In all cases, the actual values were within the 95% probability limits.

# Example: Canadian Lynx Data

A classic model fit to the log lynx data is an AR(4) (Tong, 1977).

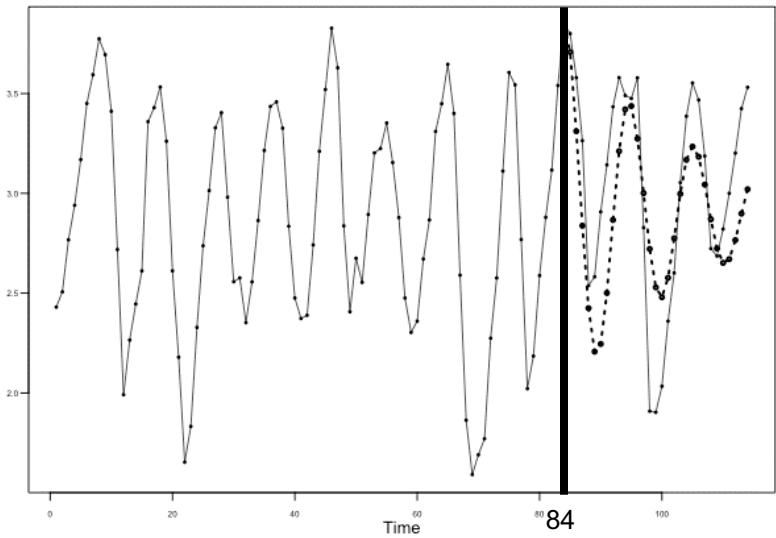
```
data(llynx)
f = fore.arma.wge(llynx, phi=c(1.3, -0.7, 0.1, 0.2), n.ahead=30,
limits=FALSE, lastn = TRUE)
```

**Forecasting last 12 values**



```
> mean((llynx[(114 - 12 + 1):114] - f$f)^2)
[1] 0.1102716
```

**Forecasting last 30 values**



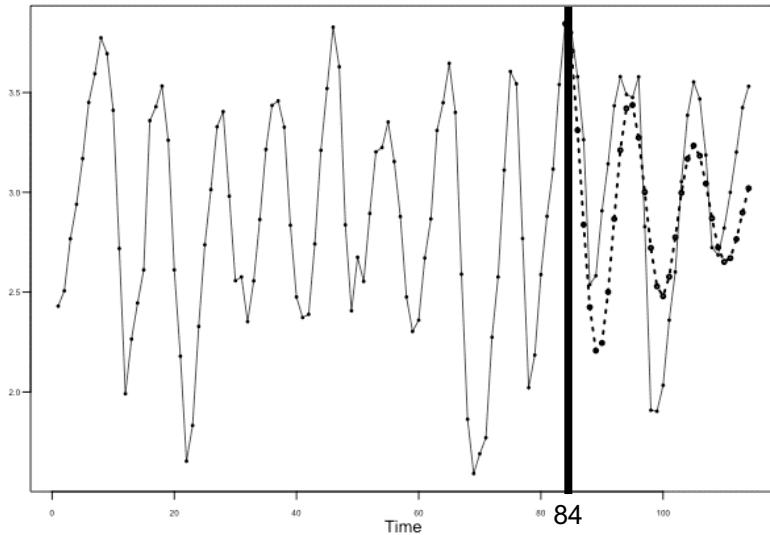
```
> mean((llynx[(114 - 30 + 1):114] - f$f)^2)
[1] 0.1440614
```

# Example: Canadian Lynx Data

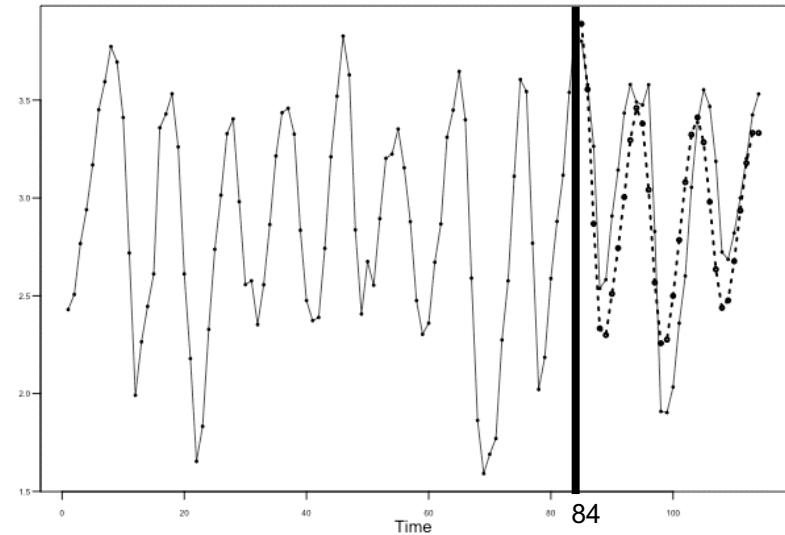
Another classic model fit to the log lynx data is an AR(11) (Tong, 1977).

```
data(llynx)
f = fore.arma.wge(llynx,phi=c(1.17, -0.54, 0.27, -0.31, 0.15, -0.15,
0.06, -0.03, 0.13, 0.20, -0.34),n.ahead=30,limits=FALSE, lastn = TRUE)
```

**Forecasting last 30 values from AR(4)**



**Forecasting last 30 values from AR(11)**



```
> mean((llynx[(114 - 30 + 1):114] - f$f)^2)
[1] 0.1440614
```

```
> mean((llynx[(114 - 30 + 1):114] - f$f)^2)
[1] 0.1003434
```

**DataScience@SMU**

# Introduction

---

## Forecasting with ARIMA

# Forecasting with ARIMA

$$(B)(1 - B)^d (X_t) = (B)a_t$$

## Comments:

- forecasts from these models are obtained using a method analogous to that used for ARMA forecasts
- we will not calculate these forecasts by hand but will typically use tswge function fore.aruma.wge
- forecast limits for these nonstationary models become unbounded as  $\ell$  increases

DataScience@SMU

# Example

---

Forecasting with ARIMA(0,1,0)

## Light board

**Example:**  $(1 - B)(X_t - \bar{X}) = a_t$

**Recall:** For stationary AR(1)—for example,  $(1 - .9B)(X_t - \mu) = a_t$

$$\hat{X}_{t_0}(\ell) = .9\hat{X}_{t_0}(\ell - 1) + \bar{X}(1 - .9)$$

**Our case:**  $(1 - B)(X_t - \bar{X}) = a_t$  i.e.  $\varphi_1 = 1$  instead of .9

$$\hat{X}_{t_0}(\ell) = \hat{X}_{t_0}(\ell - 1) + \bar{X}(1 - 1) \quad (\bar{X} \text{ plays no role in the forecasts})$$

$$\hat{X}_{t_0}(\ell) = \hat{X}_{t_0}(\ell - 1)$$

$$\hat{X}_{t_0}(1) = \hat{X}_{t_0}(0) = X_{t_0}$$

$$\hat{X}_{t_0}(2) = \hat{X}_{t_0}(1) = X_{t_0}$$

•

⋮

Continues to forecast the last data value

- Boring
- But makes sense

DataScience@SMU

# Example

---

Forecasting with ARIMA(0,1,0)

**Example:**  $(1 - B)(X_t - \bar{X}) = a_t$

**Recall:** For stationary AR(1)—for example,  $(1 - .9B)(X_t - \mu) = a_t$

$$\hat{X}_{t_0}(\ell) = .9\hat{X}_{t_0}(\ell - 1) + \bar{X}(1 - .9)$$

**Our case:**  $(1 - B)(X_t - \bar{X}) = a_t$  i.e.  $\varphi_1 = 1$  instead of .9

$$\hat{X}_{t_0}(\ell) = \hat{X}_{t_0}(\ell - 1) + \bar{X}(1 - 1) \quad (\bar{X} \text{ plays no role in the forecasts})$$

$$\hat{X}_{t_0}(\ell) = \hat{X}_{t_0}(\ell - 1)$$

$$\hat{X}_{t_0}(1) = \hat{X}_{t_0}(0) = X_{t_0}$$

$$\hat{X}_{t_0}(2) = \hat{X}_{t_0}(1) = X_{t_0}$$

•  
•  
•

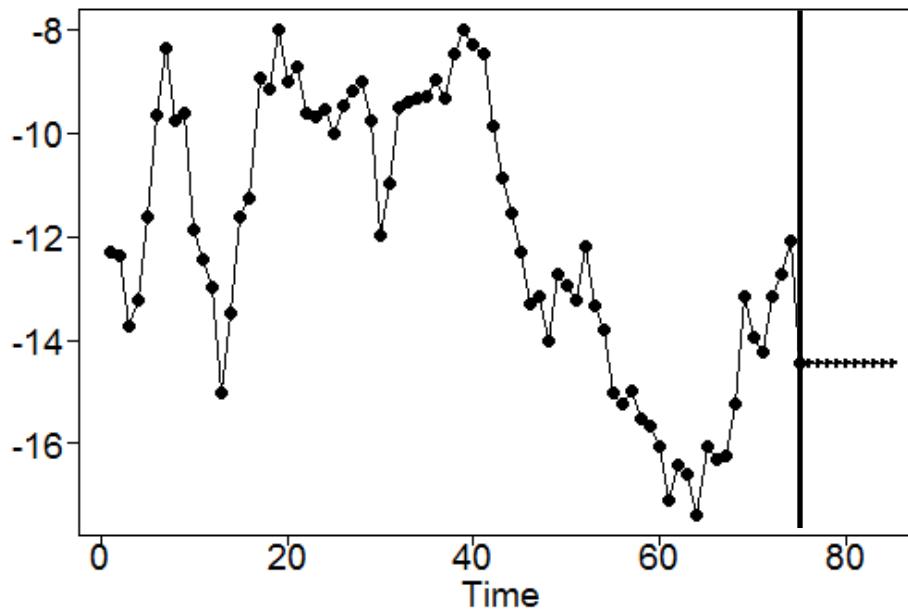
Continues to forecast the last data value

- Boring
- But makes sense

**Example:**  $(1 - B)(X_t - \bar{X}) = a_t$

`xd1=gen.aruma.wge(n=75,d=1,sn=74)`

`fore.aruma.wge(xd1,d=1,n.ahead=5,limits=F)`



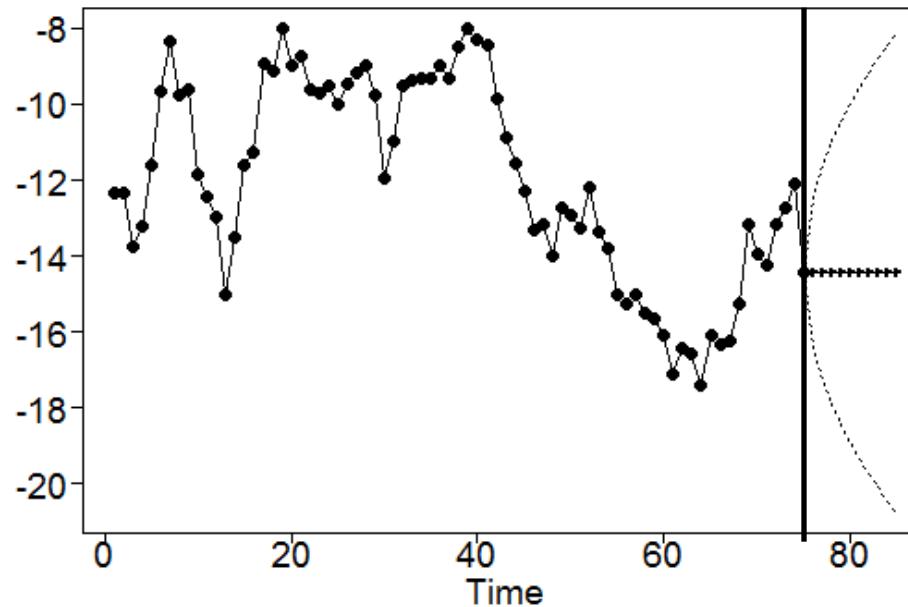
Based on the realization, it's unclear whether one should forecast the data values to go up or down.

- Forecasting the last value seems reasonable.

**Example:**  $(1 - B)(X_t - \bar{X}) = a_t$

`xd1=gen.aruma.wge(n=75,d=1,sn=74)`

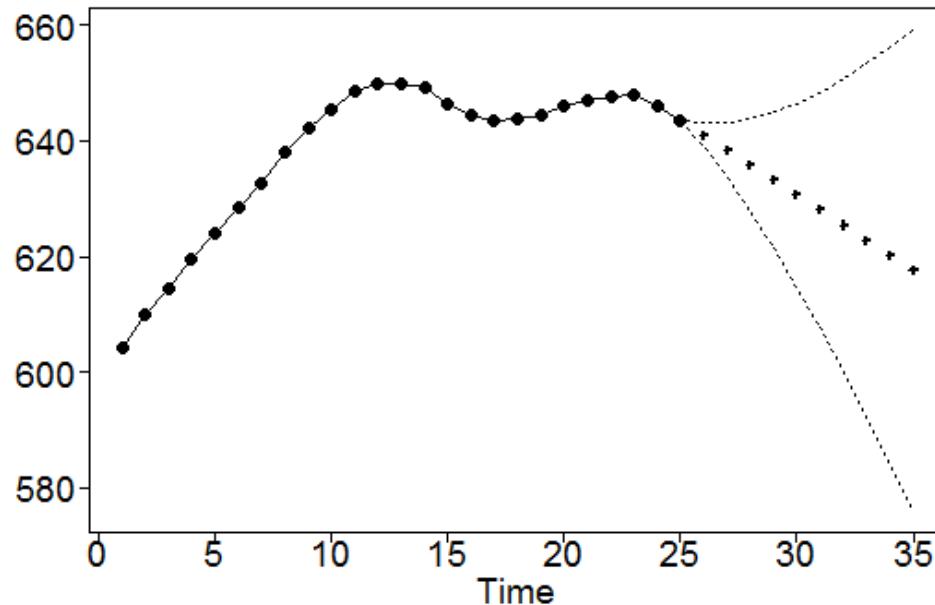
`fore.aruma.wge(xd1,d=1,n.ahead=5,limits=T)`



### Selecting 95% Probability Limits

- Shows the uncertainty associated with the forecasts
  - Could go up or down
  - Limits seem to be increasing without bound

**Example:**  $(1 - B)^2(X_t \quad ) = a_t$



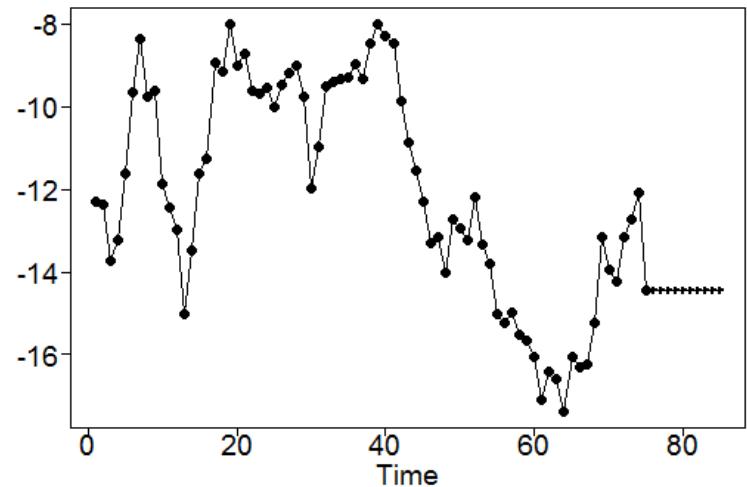
**Note:** Forecasts follow a line determined by the last 2 data values.

- That is, the most recent trend is predicted to continue.

# Key Concept!

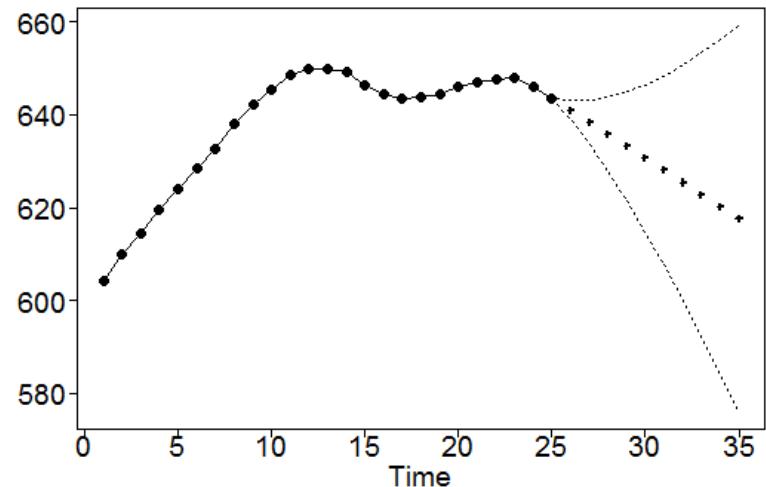
$$(1 - B)(X_t - \hat{X}_{t-1}) = a_t$$

Forecasts the last value



$$(1 - B)^2(X_t - \hat{X}_{t-2}) = a_t$$

Forecasts the trend determined  
by the last two observations



**DataScience@SMU**

tswge

---

# Generating and Forecasting Simple ARIMA Models

```
fore.aruma.wge(x,phi,theta,d,s,lambda,n.ahead,
lastn=FALSE,plot=TRUE,limits=TRUE)
```

**Example:** forecasts using the ARIMA(0,1,0) model

$$(1 - B) X_t = a_t$$

**tswge demo**

```
x=gen.aruma.wge(n=50,phi=.8,d=1,sn=15)
fore.aruma.wge(x,d=1,n.ahead=20 , limits = FALSE)
```

**Example:** forecasts using the ARIMA(1,1,0) model

$$(1 - .8B)(1 - B) X_t = a_t$$

```
x=gen.aruma.wge(n=50,phi=.8,d=1,sn=15)
fore.aruma.wge(x,phi=.8,d=1,n.ahead=20 , limits = FALSE)
```

**Example:** forecasts using the ARIMA(0,2,0) model

$$(1 - B)^2 X_t = a_t$$

```
x=gen.aruma.wge(n=50,phi=.8,d=1,sn=15)
fore.aruma.wge(x,d=2,n.ahead=20 , limits = FALSE)
```

**DataScience@SMU**

# Introduction

---

## Forecasting with Seasonal ARIMA Models

# Forecasting with Seasonal Models

$$\begin{pmatrix} 1 & B^s \end{pmatrix} (X_t) = a_t$$

**Example:**  $\begin{pmatrix} 1 & B^4 \end{pmatrix} (X_t) = a_t$

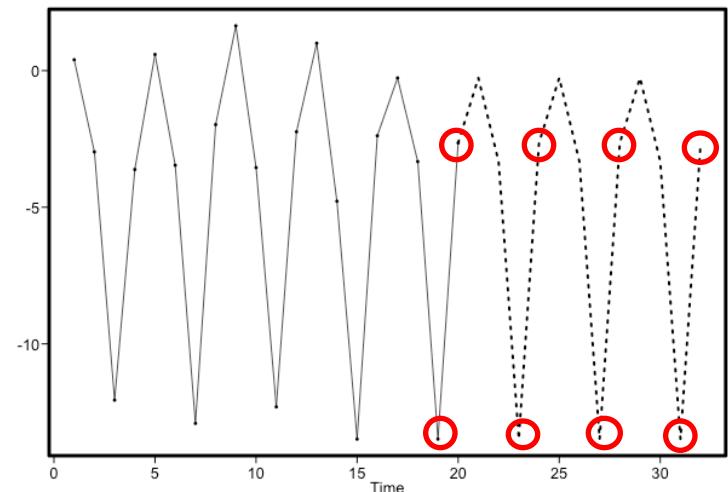
$$\hat{X}_{t_0}(\ell) = \hat{X}_{t_0}(\ell - 4) + \bar{X}(1 - 0 - 0 - 0 - 1)$$

$$\hat{X}_{t_0}(\ell) = \sum_{j=1}^p \varphi_j \hat{X}_{t_0}(\ell - j) - \sum_{j=\ell}^q \theta_j \hat{a}_{t_0+\ell-j} + \bar{X} \left[ 1 - \sum_{i=1}^p \varphi_i \right]$$

$$\hat{X}_{t_0}(\ell) = \hat{X}_{t_0}(\ell - 4)$$

## Note:

Given quarterly data, this very simple model forecasts the current quarter to be the value observed for this quarter last year.



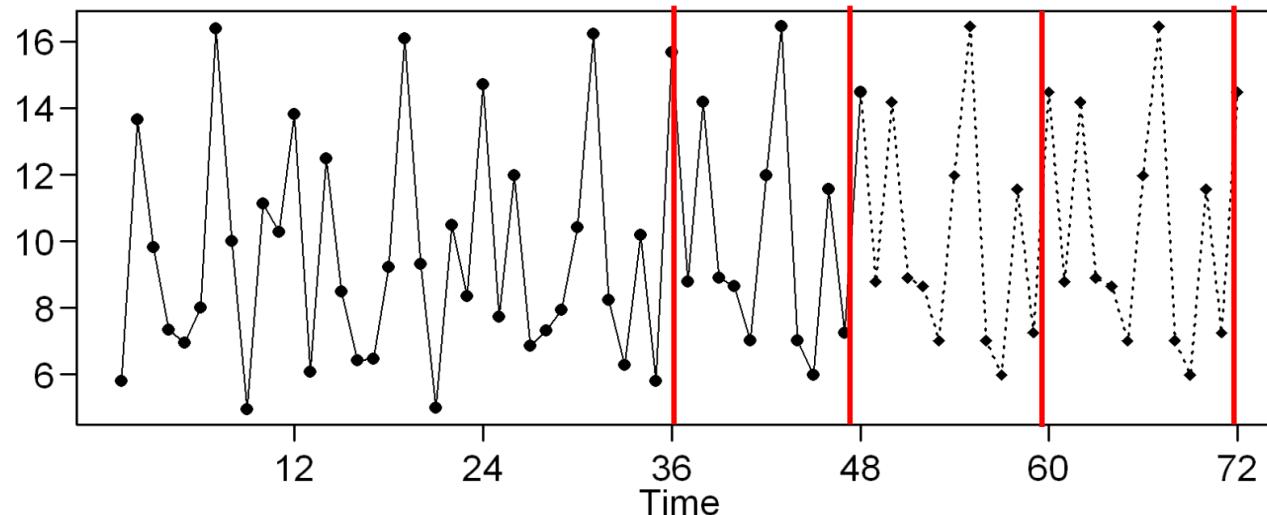
# Example:

$$\begin{pmatrix} 1 & B^s \end{pmatrix} (X_t) = a_t$$

**Forecasts:**  $\hat{X}_{t_0}(\ell) = \hat{X}_{t_0}(\ell - s)$

- That is, forecasts are an exact replica of the last  $s$  values.

Realization and forecasts for  $(1 - B^{12})(X_t) = a_t$   
as of time  $t_0 = 36$



**DataScience@SMU**

tswge

---

# Generating and Forecasting Simple Seasonal Models

## tswge demo

```
fore.aruma.wge(x,phi,theta,d,s,n.ahead,  
lastn=FALSE,plot=TRUE,limits=TRUE)
```

Forecasts from the pure seasonal model  $(1 - B^4) X_t = a_t$

```
x=gen.aruma.wge(n=20,s=4, sn = 6)  
fore.aruma.wge(x,s=4,n.ahead=8,lastn=FALSE,plot=TRUE,  
limits=FALSE)  
  
#  
x=gen.aruma.wge(n=20,s=4, sn = 6)  
fore.aruma.wge(x,s=4,n.ahead=8,lastn=TRUE,plot=TRUE,  
limits=FALSE)
```

**Example:** forecasts using seasonal model  $(1 - .8B)(1 - B^4) X_t = a_t$

```
x=gen.aruma.wge(n=20,phi=.8,s=4,sn = 6)  
fore.aruma.wge(x,phi=.8,s=4,n.ahead=8,limits=FALSE)
```

Also: lastn=TRUE

**DataScience@SMU**

# A More General Seasonal Model/ the “Airline Model”

---

# Factor Table: $s=12$

Factor table for  $(1 - B^{12})$

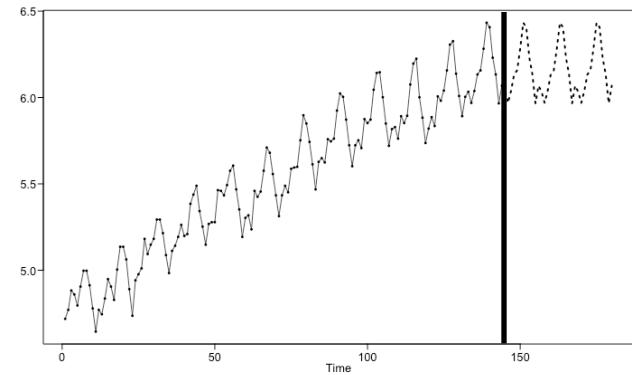
| Factor                | Abs Recip | $f$  | Root(s)         |
|-----------------------|-----------|------|-----------------|
| $1 - B$               | 1         | 0    | 1               |
| $1 - \sqrt{3}B + B^2$ | 1         | .083 | $.866 \pm .5i$  |
| $1 - B + B^2$         | 1         | .167 | $.5 \pm .866i$  |
| $1+B^2$               | 1         | .25  | $+i$            |
| $1 + B + B^2$         | 1         | .333 | $-.5 \pm .866i$ |
| $1 + \sqrt{3}B + B^2$ | 1         | .417 | $-.866 \pm .5i$ |
| $1 + B$               | 1         | .5   | -1              |

**Note:**  $(1 - B^{12})$  has a factor of  $(1 - B)$  so that the product  $(1 - B)(1 - B^{12})$  has **two** factors of  $(1 - B)$ .

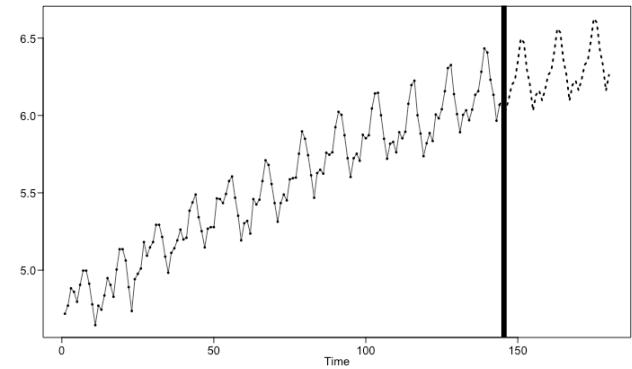
# Forecasts Using the Airline Model

- Using this model, the data for the last year are forecast into the future.
- The trend is not forecast to continue.
- In our previous examples:
  - $(1 - B) X_t = a_t$  simply forecasts the last data value into the future
  - $(1 - B)^2 X_t = a_t$  forecasts the trend to continue

$$(1 - B^{12})(X_t - \mu) = a_t$$



$$(1 - B)(1 - B^{12})(X_t - \mu) = a_t$$



## More general seasonal model

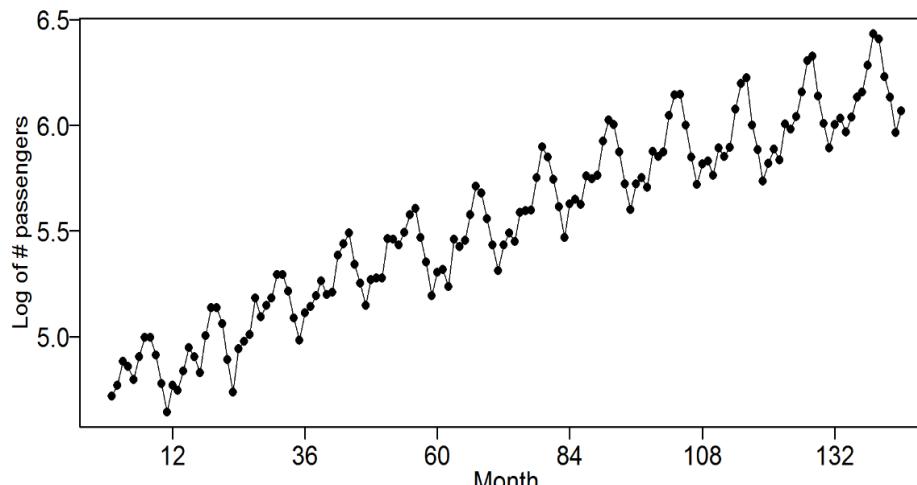
$$(B)(1 - B)(1 - B^s)(X_t \dots) = (B)a_t$$

- Forecasts follow general pattern of last  $s$  points, but not an exact replica
- “Finer detail” incorporated through  $\varphi(B)$  and  $\theta(B)$

# Example: Seasonal Model with Trend

$$(B)(1 - B)(1 - B^S)(X_t) = (B)a_t$$

- It is called the airline model.
- It is useful for modeling and forecasting data similar to the (log) airline data shown below.
- Notice that a key feature of this data set is that there is a seasonal component and a trend



**DataScience@SMU**

# Parzen, Box and Woodward / Gray Airline Models

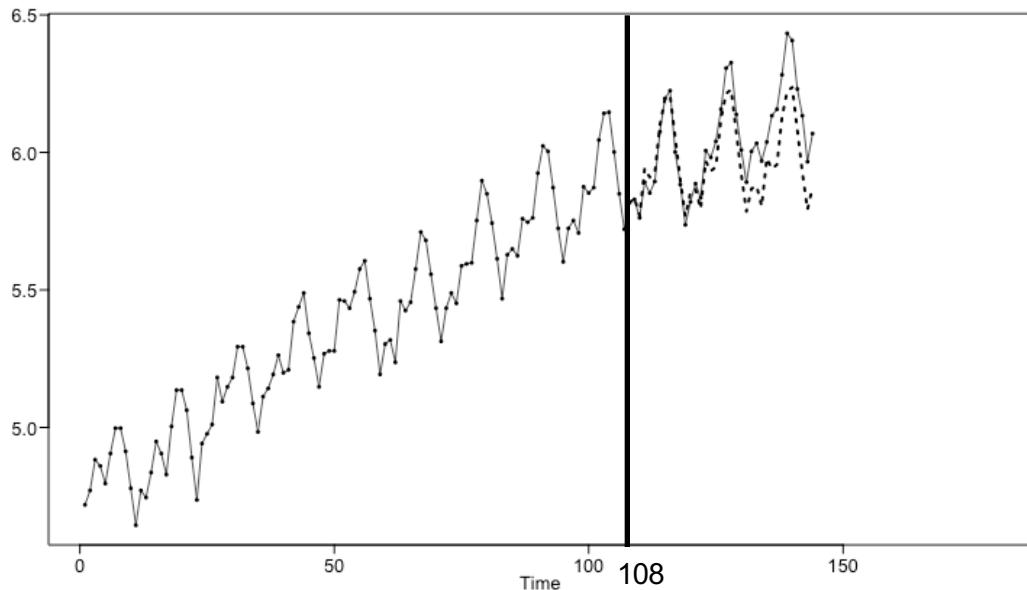
---

$$\varphi(B) (1 - B^{12})(X_t - \mu) = a_t$$

Where  $\varphi(B)$  is the 13th-order AR operator given by Parzen et al. (1980).

```
data(airlog)
```

```
fore.aruma.wge(airlog, d = 0, s = 12, phi = c(.74,0,0,0,0,0,0,0,0,.38,-.2812),n.ahead = 36,lastn = TRUE, limits = FALSE)
```



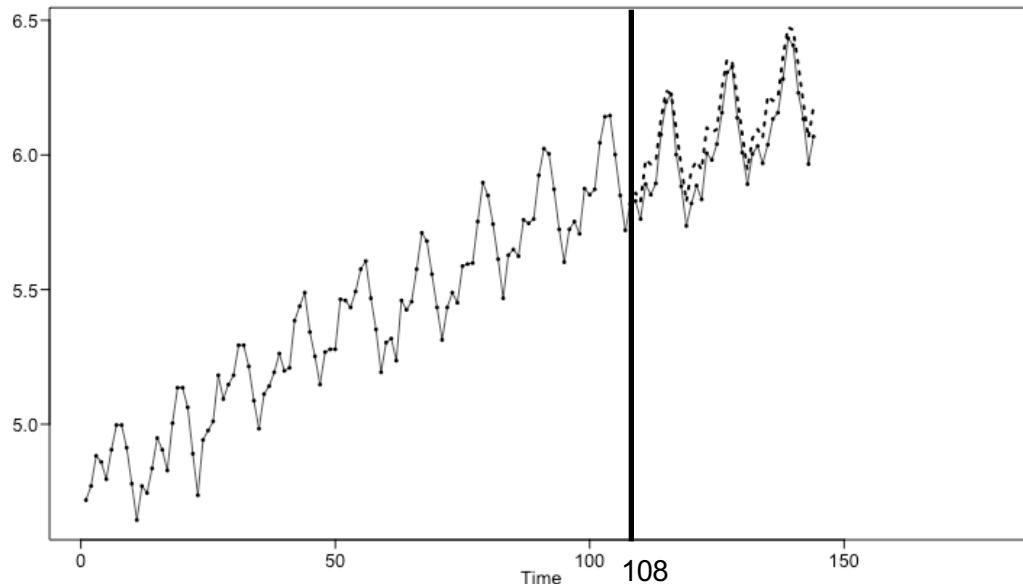
Seasonality  
predicted to  
continue... but not  
the trend

$$(1 - B)(1 - B^{12})(X_t - \mu) = \theta(B)a_t$$

Where  $\theta(B)$  is the 13th-order MA operator given by Box et al. (2008).

```
data(airlog)
```

```
fore.aruma.wge(airlog, d = 0, s = 12, phi = c(.74,0,0,0,0,0,0,0,0,.38,-.2812),n.ahead = 36,lastn = TRUE, limits = FALSE)
```

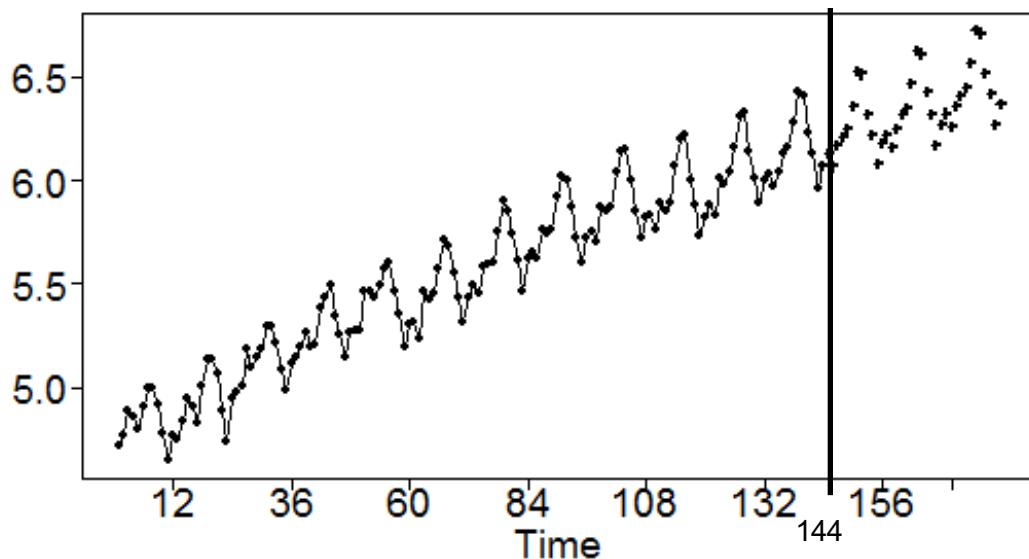


Seasonality  
predicted to  
continue... as well  
as the trend.

$$\varphi_1(B)(1 - B)(1 - B^{12}) X_t = a_t$$

Where  $\varphi_1(B)$  is the 12th-order operator given by Woodward et al. (2017).

```
data(airlog)
phi1=c(-.36,-.05,-.14,-.11,.04,.09,-.02, .02,.17,.03,-.10,-.38)
fore.aruma.wge(airlog,phi=phi1,d=1,s=12,n.ahead=36,plot=T,lastn=F,limits=F)
```

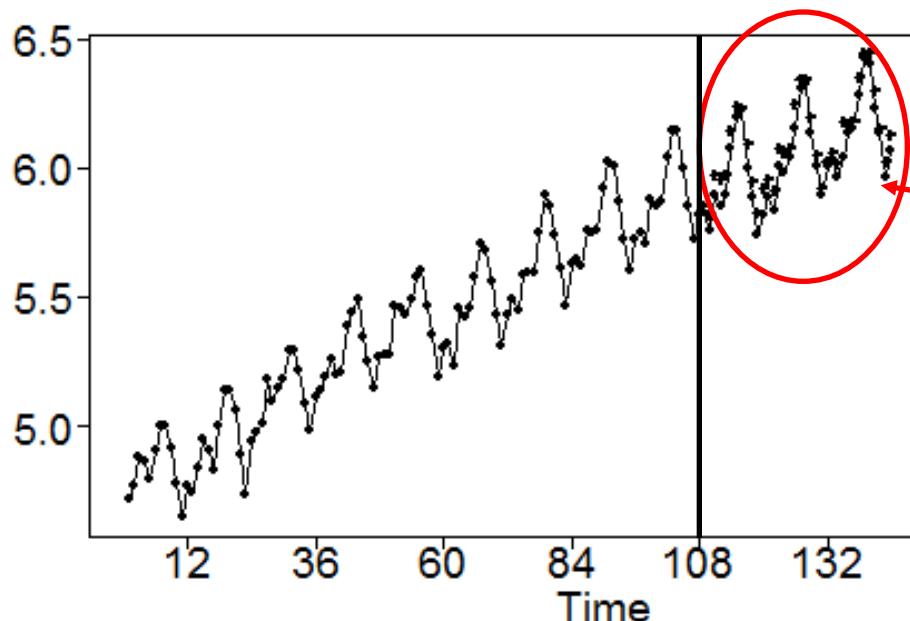


Seasonality and  
trend predicted to  
continue

$$\varphi_1(B)(1 - B)(1 - B^{12}) X_t = a_t$$

Where  $\varphi_1(B)$  is the 12th-order operator given by Woodward et al. (2017).

```
data(airlog)
phi1=c(-.36,-.05,-.14,-.11,.04,.09,-.02, .02,.17,.03,-.10,-.38)
fore.aruma.wge(airlog,phi=phi1,d=1,s=12,n.ahead=36,plot=T,lastn=T,limits=F)
```



Forecasts of last 36 months are so good that it's difficult to distinguish them from the actual values.

**DataScience@SMU**

# Competition

---

Parzen, Box, Woodward / Gray —  
Airline Forecasts, AIC, and ASE

# Airline Competition Among the Greats!

Previously, we looked at the structure of the model from three world-renowned statisticians.

```
> data(airlog)
> SA1 = artrans.wge(airlog,1) # take first differences of the data
> SA1_12 = artrans.wge(SA1,c(rep(0,11),1)) # take the 12th difference of the first difference  $(1-B)(1-B^{12})$ 
> SA12 = artrans.wge(airlog,c(rep(0,11),1)) # take the 12th difference of the data  $(1-B^{12})$ 
> Parzen = aic.wge(SA12, p = 13) # $\Phi(B)(1-B^{12})(X_t-\mu)$  = at
> Box = aic.wge(SA1_12, q = 13) #(1-B)(1-B^{12})  $(X_t-\mu)$  =  $\Theta(B)at$ 
> Woodward = aic.wge(SA1_12, p = 12) #  $\Phi(B)(1-B)(1-B^{12})(X_t-\mu)$  = at
> Parzen$value
[1] -6.465559
> Box$value
[1] -6.499275
> Woodward$value
[1] -6.423649
> # We will return to this competition when we use the model to forecast!
```

# Airline Competition Among the Greats!



$$(1 - .74B)(1 + .38B^{12})(1 - B^{12})(X_t - \mu) = a_t$$

```
> Parzen$value  
[1] -6.465559  
> Box$value  
[1] -6.499275  
> Woodward$value  
[1] -6.423649
```



$$(1 - B)(1 - B^{12})(X_t - \mu) = (1 - .4B)(1 - .6B^{12})a_t$$

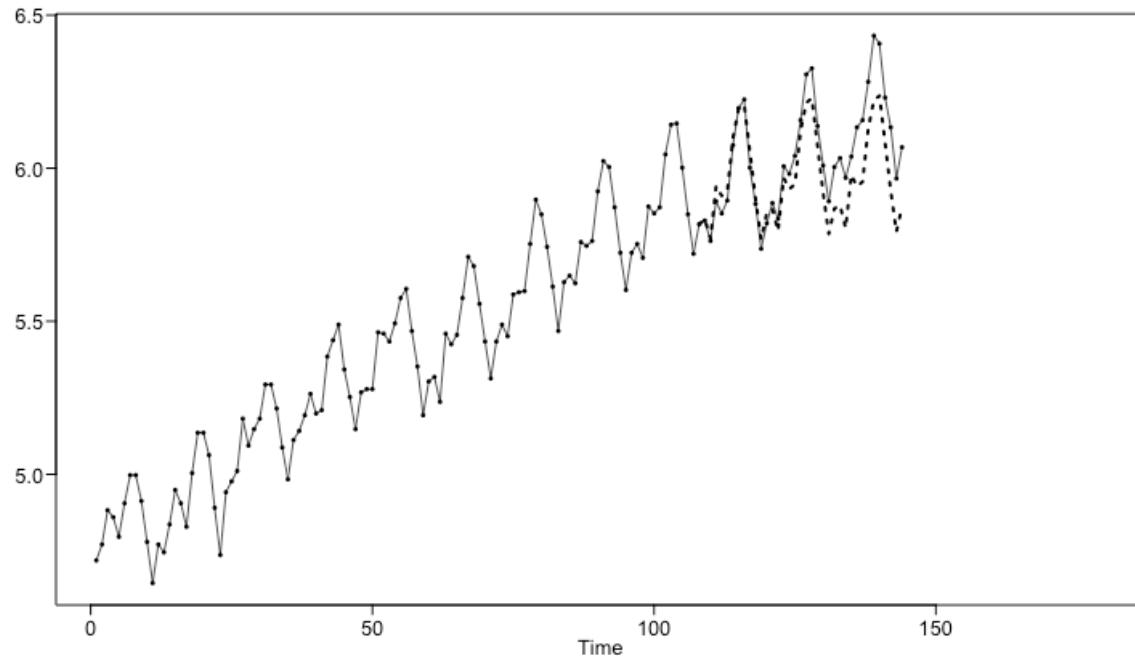


$$(1 + .36B + .05B^2 + .14B^3 + .11B^4 - .04B^5 - .09B^6 + .02B^7 - .02B^8 - .17B^9 - .03B^{10} + .10B^{11} + .38B^{12})(1 - B)(1 - B^{12})(X_t - \mu) = a_t$$

# Airline Competition Among the Greats!



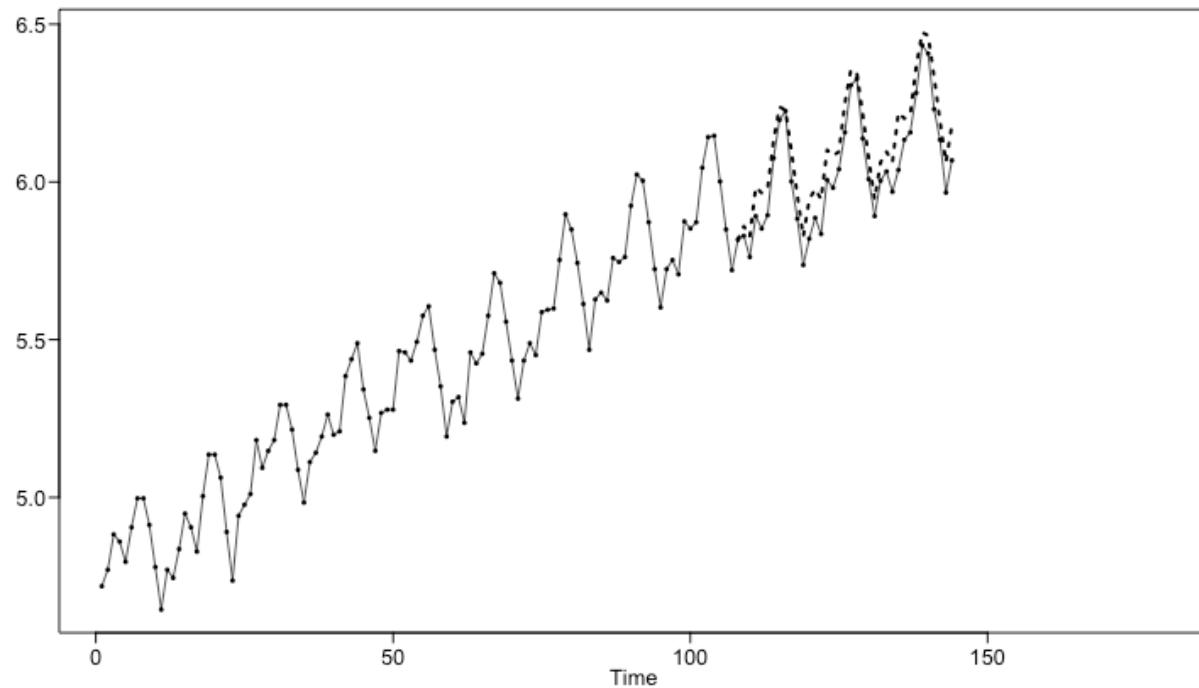
$$(1 - .74B)(1 + .38B^{12})(1 - B^{12})(X_t - \mu) = a_t$$



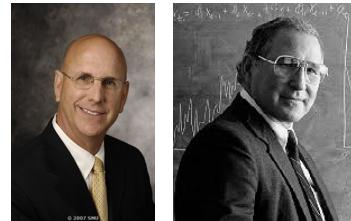
# Airline Competition Among the Greats!



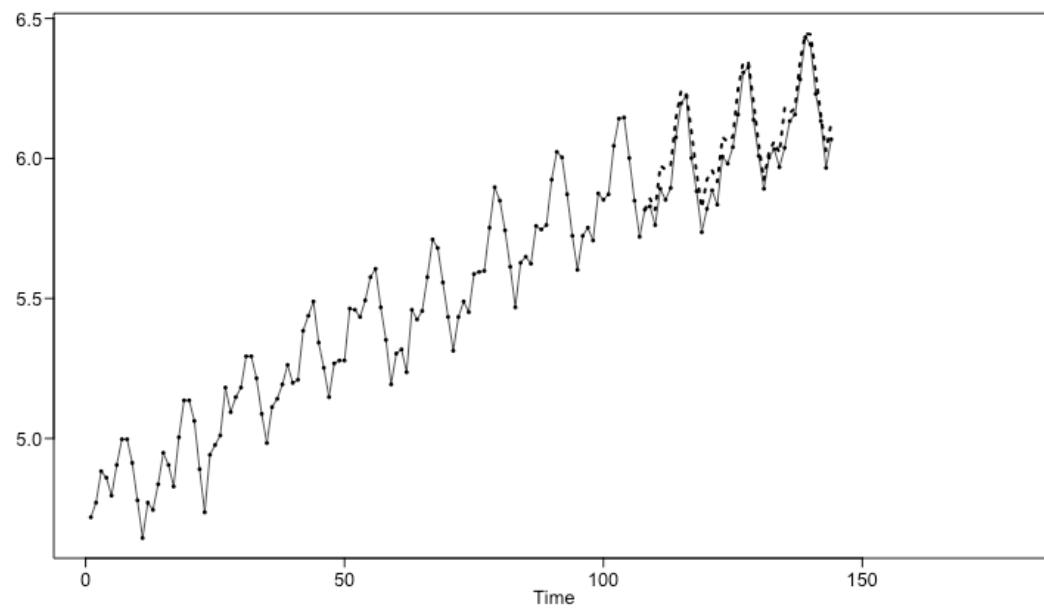
$$(I - B)(I - B^{12})(X_t - \mu) = (I - .4B)(I - .6B^{12})a_t$$



# Airline Competition Among the Greats!



$$(1 + .36B + .05B^2 + .14B^3 + .11B^4 - .04B^5 - .09B^6 + .02B^7 - .02B^8 - .17B^9 - .03B^{10} + .10B^{11} + .38B^{12})(1 - B)(1 - B^{12})(X_t - \mu) = a_t$$



# Airline Competition Among the Greats!

Now we will compare the actual models they published on the basis of the ASE calculated from forecasts of the last 36 years.



$$(1 - .74B)(1 + .38B^{12})(1 - B^{12})(X_t - \mu) = a_t$$

```
> Parzen = fore.aruma.wge(airlog, d = 0, s = 12, phi = c(.74,0,0,0,0,0,0,0,0,.38,-.2812),n.ahead = 36,lastn = TRUE, limits = FALSE)
> PARZEN_ASE = mean((airlog[(144-36+1):144] - Parzen$f)^2)
> PARZEN_ASE
[1] 0.01252636
```



$$(1 - B)(1 - B^{12})(X_t - \mu) = (1 - .4B)(1 - .6B^{12})a_t$$

```
> Box = fore.aruma.wge(airlog,d = 1, s = 12, theta = c(.4,0,0,0,0,0,0,0,0,.6,-.24),n.ahead = 36,lastn = TRUE, limits = FALSE)
> BOX_ASE = mean((airlog[(144-36+1):144] - Box$f)^2)
> BOX_ASE
[1] 0.006903242
```



$$(1 + .36B + .05B^2 + .14B^3 + .11B^4 - .04B^5 - .09B^6 + .02B^7 - .02B^8 - .17B^9 - .03B^{10} + .10B^{11} + .38B^{12})(1 - B)(1 - B^{12})(X_t - \mu) = a_t$$

?

**DataScience@SMU**

# Competition Conclusion

---

# Airline Competition Among the Greats!

Now we will compare the actual models they published on the basis of the ASE calculated from forecasts of the last 36 years.



$$(1 - .74B)(1 + .38B^{12})(1 - B^{12})(X_t - \mu) = a_t$$

AIC

```
> Parzen$value  
[1] -6.465559
```

```
> Parzen = fore.aruma.wge(airlog, d = 1, s = 12, phi = c(-.36, -.05, -.14, -.11, .04, .09, -.02, .02, .17, .03, -.10, -.38), n.ahead = 36, lastn = TRUE, limits = FALSE)  
> PARZEN_ASE = mean((airlog[(144-36+1):144] - Parzen$f)^2)  
> PARZEN_ASE  
[1] 0.01252636
```

Which model is the most useful?



$$(1 - B)(1 + .38B^{12})(1 - B^{12})(X_t - \mu) = a_t$$

```
> Box$value  
[1] -6.499275
```

```
> Box = fore.aruma.wge(airlog, d = 1, s = 12, phi = c(-.36, -.05, -.14, -.11, .04, .09, -.02, .02, .17, .03, -.10, -.38), n.ahead = 36, lastn = TRUE, limits = FALSE)  
> BOX_ASE = mean((airlog[(144-36+1):144] - Box$f)^2)  
> BOX_ASE  
[1] 0.00690324
```



$$(1 + .38B^{12})(1 - B^{12})(1 - .17B^7 + .02B^8)(X_t - \mu) = a_t$$

```
> Woodward$value  
[1] -6.423649
```

All of these professors would agree that it is nearly certain that none of these are the “right” model!

as well.

(BIC, AICc, RMSE, etc.)

$$B^7 - .02B^8$$
$$\mu) = a_t$$

```
> Woodward = fore.aruma.wge(airlog, d = 1, s = 12, phi = c(-.36, -.05, -.14, -.11, .04, .09, -.02, .02, .17, .03, -.10, -.38), n.ahead = 36, lastn = TRUE, limits = FALSE)  
> WOODWARD_ASE = mean((airlog[(144-36+1):144] - Woodward$f)^2)  
> WOODWARD_ASE  
[1] 0.004185726
```

**DataScience@SMU**

# Forecasting Signal-Plus-Noise

---

# Signal-Plus-Noise Forecasts

$$X_t = s_t + Z_t$$

**tswge allows two cases:**

$$s_t = b_0 + b_1 t$$

$$s_t = A_1 + A_2 \cos(2\pi f_0 t + A_3) = A_1 + C_1 \cos(2\pi f_0 t) + C_2 \sin(2\pi f_0 t)$$

# Forecasting Strategy (Linear Case)

Fit  $b_0$  and  $b_1$  using least squares (or use true values)

Calculate  $\hat{Z}_t = X_t - b_0 - b_1 t$

Fit AR(p) to  $\hat{Z}_t$  (covered in next unit)

Find forecasts  $\hat{Z}_{t_0}(\ell)$

Forecasts for  $X_{t_0+\ell}$  are  $\hat{X}_{t_0}(\ell) = b_0 + b_1 t + \hat{Z}_{t_0}(\ell)$

Forecast limits for  $\hat{X}_{t_0}(\ell)$  are

$$b_0 + b_1 t + \hat{Z}_{t_0}(\ell) \pm 1.96 \hat{\sigma}_a \left( \sum_{k=0}^{\ell} \frac{2}{k} \right)^2$$

where  $\hat{\sigma}_a$  and  $\hat{\sigma}_k$  are based on AR(p) model fit to  $\hat{Z}_t$

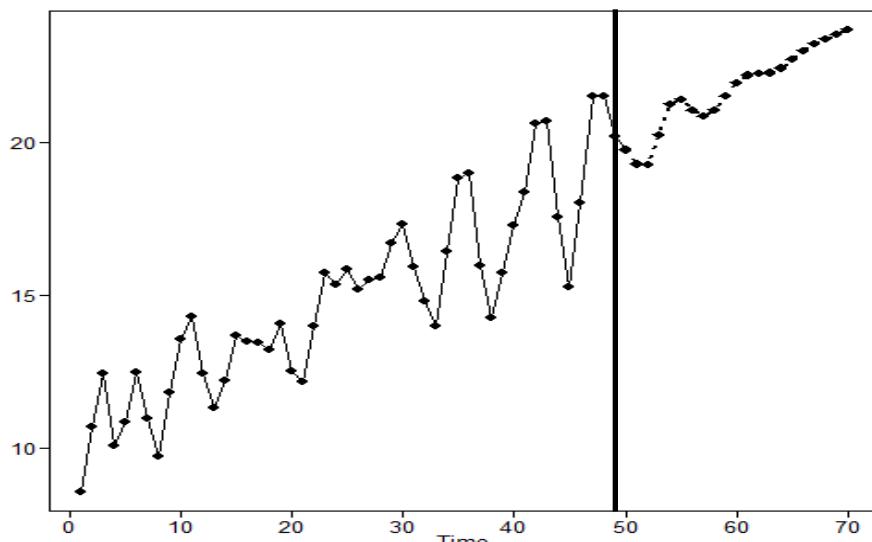
# In tswge:

```
fore.sigplusnoise.wge(x, linear = TRUE, freq = 0, max.p = 5,  
n.ahead = 10, lastn = FALSE, plot = TRUE, limits = TRUE)
```

```
x=gen.sigplusnoise.wge(n=50,b0=10,b1=.2, phi=c(.8,-.6))
```

```
#
```

```
xfore=fore.sigplusnoise.wge(x,linear=TRUE,n.ahead=20,lastn=F,limits=F)
```



Notice that the realization has a line with cyclic noise. The forecasts for early lags forecast the cyclic behavior, while for longer steps ahead, they trend to the underlying line.

# Closing Remarks

---

# Final Note

We have learned how to forecast from a **given** ARMA, ARIMA, or signal-plus-noise model.

## Important point

- The model selected makes a big difference in the forecasts.
- We saw this in the airline data: We needed a factor of  $(1 - B)$  along with  $(1 - B^{12})$  to forecast the trend to continue.
- **How do we fit the parameters to these models?**
- **How do we decide what model to use?**
  - *Upcoming Units!*

**DataScience@SMU**